

Sequential Decision Making and Reinforcement Learning

(SDMRL)

Supervised Learning vs. Model Bases Planning for Long-Term Decision-Making

Sarah Keren

The Taub Faculty of Computer Science
Technion - Israel Institute of Technology

- Model-based / model-free spectrum
- Supervised Learning for long-term decision-making
- Model-based planning for long-term decision-making

Model-free vs. Model-based Decision-making

Sequential
Decision Making
and
Reinforcement
Learning
(SDMRL)

Sarah Keren

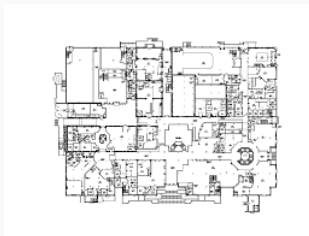
Model-free vs.
Model-based
Decision-making


Recipes for
Control
Approaches

Decision-making
as Supervised
Learning

Model-based
Planning

Planning for
Deterministic
Domains





How to come up with 'good' policies?



How to come up with 'good' policies?



How to come up with 'good' policies?
If we want to avoid collision?



How to come up with 'good' policies?

If we want to avoid collision?

If we want to turn left at the next intersection?



How to come up with 'good' policies?

If we want to avoid collision?

If we want to turn left at the next intersection?

If we want to buy apples and make it on time for dinner but have fuel for about 30 mins?

An agent typically maintains one or more of these components:

- **Model:** agent's representation of the environment
- **Value function:** how good is each state and/or action
- **Policy:** agent's behaviour function
 - **Deterministic Policy:** a mapping $\pi : \Omega \rightarrow \mathcal{A}$ from states/observations to actions.
 - **Stochastic Policy:** a mapping $\pi : \mathcal{S} \times \mathcal{A} \rightarrow [0, 1]$ from state and action pairs to the probability $\pi(a|s)$ of taking action a when in state s .

- Expected return for episodic tasks:

$$G_t = R_{t+1} + R_{t+2} + R_{t+3} + \dots + R_T = \sum_{k=0}^T R_{t+k+1}$$

- Expected return for continuing tasks:

$$G_t = R_{t+1} + \gamma R_{t+2} + \gamma^2 R_{t+3} + \dots = \sum_{k=0}^{\infty} \gamma^k R_{t+k+1}$$

when γ is the discount factor.

- Returns at successive time steps are related to each other:

$$G_t = R_{t+1} + \gamma(R_{t+2} + \gamma R_{t+3} + \gamma^2 R_{t+4} + \dots) = R_{t+1} + \gamma G_{t+1}$$

Value Functions

Value functions estimate how ‘good’ it is for the agent to be in a given state or how good it is to perform a given action in a given state.

“How good” is defined in terms of expected return.

Since the rewards the agent can expect to receive in the future depend on what actions it will take, value functions are defined with respect to particular policies.

State-Value Function for Policy π :

$$v_{\pi}(s) \doteq \mathbb{E}_{\pi} [G_t \mid S_t = s]$$

Action-Value Function for Policy π :

$$q_{\pi}(s, a) \doteq \mathbb{E}_{\pi} [G_t \mid S_t = s, A_t = a]$$

What is G_t ?

Reinforcement
Learning

(SDMRL)

Sarah Keren

Model-free vs.
Model-based
Decision-making

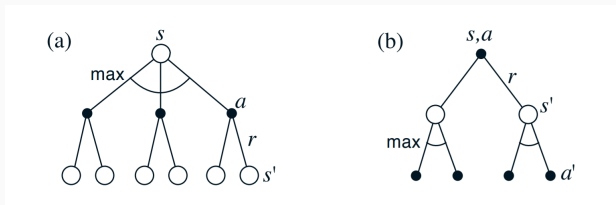
Recipes for
Control
Approaches

Decision-making
as Supervised
Learning

Model-based
Planning

Planning for
Deterministic
Domains

Value Functions



Backup diagrams for $v_\pi(s)$ (left) and $q_\pi(s, a)$ (right).


Since it may not be practical to keep separate averages for each state individually v_π and q_π are often represented as parameterized functions (with fewer parameters than states).

How to find an optimal policy ?

Recipes for Control Approaches

Policy Evaluation and Policy Update

- Prediction / Evaluation: evaluate the future given a policy
- Control/ Update / Improvement: optimize the policy.



Control Approaches Skeleton

Reinforcement
Learning

(SDMRL)

Sarah Keren



Model-free vs.
Model-based
Decision-making

Recipes for
Control
Approaches

Decision-making
as Supervised
Learning

Model-based
Planning

Planning for
Deterministic
Domains



Left image by Sutton and Barto. Right image By Sergey Levine

Model Based vs. Model Free

- For this distinction, a **model** typically refers to the transition function \mathcal{P} and the reward function \mathcal{R} .
- A model free approach only maintains a policy or value function, but no model.
- A model-based approach maintains a policy or value function and a model.



https://www.davidsilver.uk/wp-content/uploads/2020/03/intro_RL.pdf

Value vs. Policy Based

- **Value Based:**
 - No Policy (Implicit)
 - Value Function
- **Policy Based**
 - Policy
 - No Value Function
- **Combined approach (a.k.a Actor-Critic)**
 - Policy
 - Value Function

Reinforcement
Learning

(SDMRL)

Sarah Keren

Model-free vs.
Model-based
Decision-making

Recipes for
Control
Approaches

Decision-making
as Supervised
Learning

Model-based
Planning

Planning for
Deterministic
Domains

On Policy vs. Off Policy

According to Sutton and Barto 2018

- **On-policy** methods evaluate and improve the policy based on the policy that is used to make decisions.
- **Off-policy** methods evaluate or improve a policy different from that used to generate the data.
- We distinguish between the **behavior policy**, according to which an agent interacts with the environment and **target policy** the policy the agent is trying to learn that will optimize its utility.
- In on-policy methods behavior policy == target policy. In off-policy methods they are different.

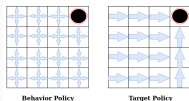


Image from <https://towardsdatascience.com/>

Reinforcement
Learning

(SDMRL)

Sarah Keren

Model-free vs.
Model-based
Decision-making

Recipes for
Control
Approaches


Decision-making
as Supervised
Learning

Model-based
Planning

Planning for
Deterministic
Domains

Online vs. Offline

- **Online** iteratively collecting data (a.k.a. experiences) by interacting with the environment.
- **Offline** utilize previously collected data, without additional online data collection.
 - resembles the standard supervised learning.
 - make it possible to turn large datasets into powerful decision making engines.
 - Batch Reinforcement Learning, behavioral cloning



From Or Rivlin <https://towardsdatascience.com/the-power-of-offline-reinforcement-learning-5e3d3942421c>

Planning and Replanning

- **Planning:** the task of coming up with a sequence of actions that will achieve a goal.
 - **Sensorless (conformant) planning:** constructing sequential plans to be executed without perception
 - **Conditional (contingent) planning:** constructing a conditional plan with different branches for different contingencies that could happen.
 - **Continuous planning:** a planner designed to persist over a lifetime.
- **Execution monitoring and replanning:** constructing a plan, but monitoring its execution and generating a new plan when necessary.

from <https://www.cpp.edu/~ftang/courses/CS420/notes/planning.pdf>

Reinforcement
Learning

(SDMRL)

Sarah Keren

Model-free vs.
Model-based
Decision-making

Recipes for
Control
Approaches

Decision-making
as Supervised
Learning

Model-based
Planning

Planning for
Deterministic
Domains

Solution Approaches




image by David Silver

Reinforcement
Learning

(SDMRL)

Sarah Keren

Model-free vs.
Model-based
Decision-making

Recipes for
Control
Approaches

Decision-making
as Supervised
Learning

Model-based
Planning

Planning for
Deterministic
Domains

Approaches to Control

- Supervised learning
- Model-Based Planning
- Monte-Carlo methods
- Temporal-Difference methods
- Combined approaches

Model Free

Model Based

Reinforcement
Learning

(SDMRL)

Sarah Keren

Model-free vs.
Model-based
Decision-making

Recipes for
Control
Approaches

Decision-making
as Supervised
Learning

Model-based
Planning

Planning for
Deterministic
Domains

Decision-making as Supervised Learning

Behavior Cloning: Example

Reinforcement
Learning

(SDMRL)

Sarah Keren


Model-free vs.
Model-based
Decision-making

Recipes for
Control
Approaches

Decision-making
as Supervised
Learning

Model-based
Planning

Planning for
Deterministic
Domains



- The objective is to learn a policy by mimicking the behavior demonstrated by an expert.
- A model (policy) $\pi_{\theta}(s)$ is trained to minimize the discrepancy between its predicted actions and the expert's actions over the collected dataset.

Can be broken down into the following steps:

- **Data Collection:** Gather a dataset of state-action pairs (s_i, a_i) , where s_i represents the state and a_i represents the action taken by the expert at that state.
- **Model Selection:** Choose a model to approximate the expert's policy. Typically, a neural network $\pi_{\theta}(s)$ parameterized by θ .

Ideas for managing the learning process?

Behavioral cloning

- **Loss Function:** measures the difference between the expert's actions and the actions predicted by the model. For example, mean squared error (MSE) or cross-entropy loss for discrete actions.

$$L(\theta) = \frac{1}{N} \sum_{i=1}^N \|a_i - \pi_{\theta}(s_i)\|^2$$

where N is the number of state-action pairs in the dataset.

- **Training:** Optimize the model parameters θ by minimizing the loss function (e.g., using gradient descent)

$$\theta^* = \arg \min_{\theta} L(\theta)$$

- **Policy Execution:** Use the trained model $\pi_{\theta^*}(s)$ as the policy to predict actions for new states during execution.

Limitations?

Reinforcement
Learning

(SDMRL)

Sarah Keren

Model-free vs.
Model-based
Decision-making


Recipes for
Control
Approaches

Decision-making
as Supervised
Learning

Model-based
Planning

Planning for
Deterministic
Domains

Dagger (Dataset Aggregation)

- 
1. train $\pi_\theta(\mathbf{a}_t|\mathbf{o}_t)$ from human data $\mathcal{D} = \{\mathbf{o}_1, \mathbf{a}_1, \dots, \mathbf{o}_N, \mathbf{a}_N\}$
 2. run $\pi_\theta(\mathbf{a}_t|\mathbf{o}_t)$ to get dataset $\mathcal{D}_\pi = \{\mathbf{o}_1, \dots, \mathbf{o}_M\}$
 3. Ask human to label \mathcal{D}_π with actions \mathbf{a}_t
 4. Aggregate: $\mathcal{D} \leftarrow \mathcal{D} \cup \mathcal{D}_\pi$

<https://jonathan-hui.medium.com/rl-imitation-learning-ac28116c02fc>

<https://bair.berkeley.edu/blog/2017/10/26/dart/>

Reinforcement
Learning

(SDMRL)

Sarah Keren

Model-free vs.
Model-based
Decision-making

Recipes for
Control
Approaches

Decision-making
as Supervised
Learning

Model-based
Planning

Planning for
Deterministic
Domains

Model-based Planning

Problem-solving agents

Restricted form of general agent

```
def Simple-Problem-Solving-Agent (problem):  
    state, some description of the current world state  
    seq, an action sequence, initially empty  
    state ← UPDATE-STATE(state, percept)  
    if seq is empty then  
        seq ← SEARCHFOR SOLUTION(problem)  
    action ← SELECT ACTION(seq, state)  
    seq ← REMAINDER(seq, action)  
    return action
```

- Offline problem solving; solution executed “eyes closed.”
- Based on a model of the environment and its dynamics

Reinforcement
Learning

(SDMRL)

Sarah Keren

Model-free vs.
Model-based
Decision-making


Recipes for
Control
Approaches

Decision-making
as Supervised
Learning

Model-based
Planning

Planning for
Deterministic
Domains

Examples



Reinforcement Learning

(SDMRL)

Sarah Keren

Model-free vs.
Model-based
Decision-making

Recipes for
Control
Approaches

Decision-making
as Supervised
Learning

Model-based
Planning

Planning for
Deterministic
Domains

Model-Based Planning

Characterized by:

- A set of **states** \mathcal{S} a system can be in.
 - a state is a full assignment to the set of **variables** (features) \mathcal{X} .
- **Actions** \mathcal{A} change the values of certain variables.
- **Reward Function** \mathcal{R} sets a numeric signal passed from the environment (can represent cost)
 - used to signal the objective
 - some domains have **goal conditions** such that the agent should reach **goal states** that satisfy is(e.g., 'be at Austin').
- **Objective:** find a **policy** that drives the initial state into a goal state or that maximizes the expected accumulated reward.
- Language is **generic** and not domain specific.
- **Complexity:** Even in the simplest setting it is NP-hard; i.e., exponential in the number of variables in the worst case.

Reinforcement
Learning

(SDMRL)

Sarah Keren

Model-free vs.
Model-based
Decision-making

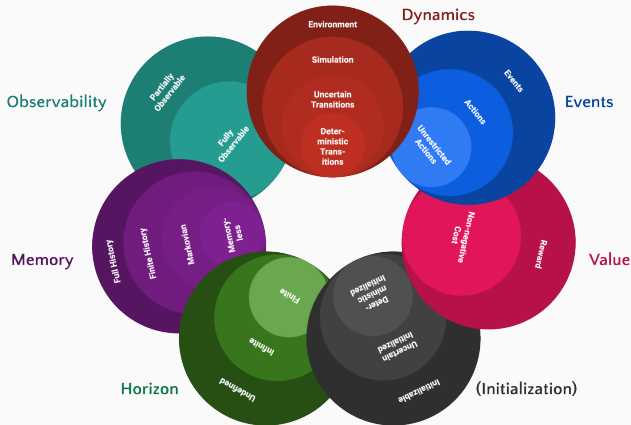
Recipes for
Control
Approaches

Decision-making
as Supervised
Learning

Model-based
Planning

Planning for
Deterministic
Domains

Many Forms of Planning



[https://github.com/fteicht/
icaps24-skdecide-tutorial/tree/main](https://github.com/fteicht/icaps24-skdecide-tutorial/tree/main)

Reinforcement
Learning

(SDMRL)

Sarah Keren

Model-free vs.
Model-based
Decision-making

Recipes for
Control
Approaches

Decision-making
as Supervised
Learning

Model-based
Planning

Planning for
Deterministic
Domains

Planning for Deterministic Domains

Planning and sequential decision making

In the classical planning setting:

Plans (aka **solutions**) are sequences of moves that transform the initial state into the goal state

What is our task?

- 1 Find out whether there is a solution
- 2 Find any solution
- 3 Find an optimal solution
- 4 Find near-optimal solution
- 5 Fixed amount of time, find best solution possible
- 6 Find solution that satisfy property \mathbb{N}
(what is \mathbb{N} ? you choose!)

♠ While all these tasks sound related, they are **very different**. The best suited techniques are almost disjoint.

Reinforcement
Learning

(SDMRL)

Sarah Keren

Model-free vs.
Model-based
Decision-making

Recipes for
Control
Approaches

Decision-making
as Supervised
Learning

Model-based
Planning

Planning for
Deterministic
Domains

Heuristic Search

Local search

Heuristics

Back to deterministic transition systems

A transition system is **deterministic** if there is only **one initial state** and all **actions are deterministic**. Hence all future states of the world are completely predictable.

Definition (deterministic transition system)

A **deterministic transition system** is $\langle S, s_0, A, S_G, c \rangle$ where

- finite state space S
- an initial state $s_0 \in S$
- a set $S_G \subseteq S$ of goal states
- applicable actions $A(s) \subseteq A$ for $s \in S$
- a transition function $s' = f(a, s)$ for $a \in A(s)$
- a cost function $c : A^* \rightarrow [0, \infty)$

Reinforcement
Learning

(SDMRL)

Sarah Keren

Model-free vs.
Model-based
Decision-making

Recipes for
Control
Approaches

Decision-making
as Supervised
Learning

Model-based
Planning

Planning for
Deterministic
Domains


Heuristic Search

Local search

Heuristics

Blocks World

blocks	states
1	1
2	
3	1
4	7
5	50
6	405
7	3763
8	39435
9	459655
...	
19	13564373693588558173



- 1 Finding a solution is polynomial time in the number of blocks (move everything onto the table and then construct the goal configuration).
- 2 Finding a shortest solution is NP-hard ...

Planning problems (reminder)

- Route selection (from Arad to Bucharest)
- Solving 15-puzzle (or Rubik's cube, or ...)
- Selecting and ordering movements of an elevator or a crane
- Production lines control
- Autonomous robots
- Crisis management
- ...

Reinforcement
Learning

(SDMRL)

Sarah Keren

Model-free vs.
Model-based
Decision-making

Recipes for
Control
Approaches

Decision-making
as Supervised
Learning

Model-based
Planning

Planning for
Deterministic
Domains

Heuristic Search

Local search

Heuristics


- **State-space search**: one of the big success stories of AI
- Must carefully distinguish two different problems:
 - **satisficing planning**: any solution is OK (although shorter solutions typically preferred)
 - **optimal planning**: plans must have shortest possible length
- Both are often solved by search, but:
 - details are **very different**
 - almost **no overlap** between good techniques for satisficing planning and good techniques for optimal planning
 - many problems that are trivial for satisficing planners are impossibly hard for optimal planners

Planning by forward search: progression

Progression: Computing the successor state $f(s, a)$ of a state s with respect to an action a .

Progression planners find solutions by forward search:

- start from initial state
- iteratively pick a previously generated state and **progress it** through an action, generating a new state
- solution found when a goal state generated




Planning by forward search: progression

Progression: Computing the successor state $f(s, a)$ of a state s with respect to an action a .

Progression planners find solutions by forward search:

- start from initial state
- iteratively pick a previously generated state and **progress it** through an action, generating a new state
- solution found when a goal state generated




Planning by forward search: progression

Progression: Computing the successor state $f(s, a)$ of a state s with respect to an action a .

Progression planners find solutions by forward search:

- start from initial state
- iteratively pick a previously generated state and **progress it** through an action, generating a new state
- solution found when a goal state generated



Classification of search algorithms

uninformed search vs. heuristic search:

- **uninformed search algorithms** only use the basic ingredients for general search algorithms
- **heuristic search algorithms** additionally use **heuristic functions** which estimate how close a node is to the goal

systematic search vs. local search:

- **systematic algorithms** consider a large number of search nodes simultaneously
- **local search algorithms** work with one (or a few) candidate solutions (search nodes) at a time
- not a black-and-white distinction; there are **crossbreeds** (e.g., enforced hill-climbing)

Reinforcement
Learning

(SDMRL)

Sarah Keren

Model-free vs.
Model-based
Decision-making

Recipes for
Control
Approaches

Decision-making
as Supervised
Learning

Model-based
Planning

Planning for
Deterministic
Domains

Heuristic Search

Local search

Heuristics

Popular uninformed systematic search algorithms:

- breadth-first search
- depth-first search
- iterated depth-first search

Popular uninformed local search algorithms:

- random walk

Evaluating Algorithms

Dimensions for evaluation

- **completeness**: always find a solution if one exists?
- **time complexity**: number of nodes generated/expanded
- **space complexity**: number of nodes in memory
- **optimality**: does it always find a least-cost solution?

Time/space complexity measured in terms of

- b* maximum branching factor of the search tree
- d* depth of the least-cost solution
- m* maximum depth of the state space (may be ∞)

Reinforcement
Learning

(SDMRL)

Sarah Keren

Model-free vs.
Model-based
Decision-making

Recipes for
Control
Approaches

Decision-making
as Supervised
Learning

Model-based
Planning

Planning for
Deterministic
Domains

Heuristic Search

Local search

Heuristics

Properties of breadth-first search

Complete Yes (if b is finite)

Time $1 + b + b^2 + b^3 + \dots + b^d + b(b^d - 1) = O(b^{d+1})$, i.e.,
 $\exp(d)$

Space $O(b^{d+1})$ (*why?*)

Optimal Yes (if cost = 1 per step); can be generalized (*how?*)

Space is the big problem; can easily generate nodes at
100MB/sec so 24hrs = 8640GB.

Reinforcement
Learning

(SDMRL)

Sarah Keren

Model-free vs.
Model-based
Decision-making

Recipes for
Control
Approaches

Decision-making
as Supervised
Learning

Model-based
Planning

Planning for
Deterministic
Domains

Heuristic Search

Local search

Heuristics

Heuristic search algorithms: systematic

- Heuristic search algorithms are the most common and overall most successful algorithms for deterministic planning.

Popular systematic heuristic search algorithms:

- greedy best-first search
- A^*
- weighted A^*
- id- a^*
- depth-first branch-and-bound search
- breadth-first heuristic search
- ...

Reinforcement
Learning

(SDMRL)

Sarah Keren

Model-free vs.
Model-based
Decision-making

Recipes for
Control
Approaches

Decision-making
as Supervised
Learning

Model-based
Planning

Planning for
Deterministic
Domains

Heuristic Search

Local search

Heuristics

Heuristic search algorithms: local

- Heuristic search algorithms are the most common and overall most successful algorithms for deterministic planning.

Popular heuristic local search algorithms:

- hill-climbing
- enforced hill-climbing
- beam search
- tabu search
- genetic algorithms
- simulated annealing
- ...

Reinforcement
Learning

(SDMRL)

Sarah Keren

Model-free vs.
Model-based
Decision-making

Recipes for
Control
Approaches

Decision-making
as Supervised
Learning

Model-based
Planning


Planning for
Deterministic
Domains

Heuristic Search

Local search

Heuristics

Heuristic search: idea



Reinforcement
Learning

(SDMRL)

Sarah Keren

Model-free vs.
Model-based
Decision-making

Recipes for
Control
Approaches

Decision-making
as Supervised
Learning

Model-based
Planning

Planning for
Deterministic
Domains

Heuristic Search

Local search

Heuristics

Required ingredients for heuristic search

A **heuristic search algorithm** requires one more operation in addition to the definition of a search space.

Definition (heuristic function)

Let Σ be the set of nodes of a given search space.

A **heuristic function** or **heuristic** (for that search space) is a function $h : \Sigma \rightarrow \mathbb{N}_0 \cup \{\infty\}$.

- The value $h(\sigma)$ supposed to estimate the distance from node σ to the nearest goal node.
- Typically: $h(\sigma) \stackrel{\text{def}}{=} h(\text{state}(\sigma))$

Reinforcement
Learning

(SDMRL)

Sarah Keren

Model-free vs.
Model-based
Decision-making

Recipes for
Control
Approaches

Decision-making
as Supervised
Learning

Model-based
Planning

Planning for
Deterministic
Domains

Heuristic Search

Local search

Heuristics

What exactly is a heuristic estimate?

What does it mean that h “estimates the goal distance”?

- For most heuristic search algorithms, h does not need to have any strong properties for the algorithm to work
- However, the **efficiency** of the algorithm closely relates to how accurately h reflects the actual goal distance.
- For some algorithms, like A*, we can prove strong formal relationships between properties of h and properties of the algorithm (optimality, dominance, run-time for bounded error, ...)
- For other search algorithms, “it works well in practice” is often as good an analysis as one gets.

Reinforcement
Learning

(SDMRL)

Sarah Keren

Model-free vs.
Model-based
Decision-making

Recipes for
Control
Approaches

Decision-making
as Supervised
Learning

Model-based
Planning

Planning for
Deterministic
Domains

Heuristic Search

Local search

Heuristics

Let Σ be the set of nodes of a given search space.

Definition (optimal/perfect heuristic)

The **optimal** or **perfect heuristic** of a search space is the heuristic h^* which maps each search node σ to the length of a shortest path from $state(\sigma)$ to any goal state.

Note: $h^*(\sigma) = \infty$ iff no goal state is reachable from σ .

A heuristic h is called

- **safe** if for all $\sigma \in \Sigma$ $h(\sigma) = \infty$ implies $h^*(\sigma) = \infty$
- **goal-aware** if $h(\sigma) = 0$ for all goal nodes $\sigma \in \Sigma$
- **admissible** if $h(\sigma) \leq h^*(\sigma)$ for all nodes $\sigma \in \Sigma$
- **consistent** if $h(\sigma) \leq h(\sigma') + cost(\sigma, \sigma')$
for all nodes $\sigma, \sigma' \in \Sigma$ such that σ' is a successor of σ

Relationships?

Greedy best-first search

Greedy best-first search (with duplicate detection)

```
open := new min-heap ordered by  $(\sigma \mapsto h(\sigma))$ 
open.insert(make-root-node(init()))
closed :=  $\emptyset$ 
while not open.empty():
     $\sigma = open.pop\_min()$ 
    if state( $\sigma$ )  $\notin$  closed:
        closed := closed  $\cup$  {state( $\sigma$ )}
        if is-goal(state( $\sigma$ )):
            return extract-solution( $\sigma$ )
        for each  $\langle o, s \rangle \in succ(state(\sigma))$ :
             $\sigma' := make\_node(\sigma, o, s)$ 
            if  $h(\sigma') < \infty$ :
                open.insert( $\sigma'$ )
return unsolvable
```

Reinforcement
Learning

(SDMRL)

Sarah Keren

Model-free vs.
Model-based
Decision-making

Recipes for
Control
Approaches

Decision-making
as Supervised
Learning

Model-based
Planning

Planning for
Deterministic
Domains

Heuristic Search

Local search

Heuristics

Planning trip to Bucharest with SLD heuristic

Arad	366	Mehadia	241
Bucharest	0	Neamt	234
Craiova	160	Oradea	380
Dobreta	242	Pitesti	100
Eforie	161	Rimnicu Vilcea	193
Fagaras	176	Sibiu	253
Giurgiu	77	Timisoara	329
Hirsova	151	Urziceni	80
Iasi	226	Vaslui	199
Lugoj	244	Zerind	374

Reinforcement
Learning

(SDMRL)

Sarah Keren

Model-free vs.
Model-based
Decision-making

Recipes for
Control
Approaches

Decision-making
as Supervised
Learning

Model-based
Planning

Planning for
Deterministic
Domains

Heuristic Search
Local search
Heuristics

GBFS by example


(a) The initial state




(b) After expanding Arad



(c) After expanding Sibiu



(d) After expanding Fagaras



Reinforcement
Learning

(SDMRL)

Sarah Keren

Model-free vs.
Model-based
Decision-making

Recipes for
Control
Approaches

Decision-making
as Supervised
Learning

Model-based
Planning

Planning for
Deterministic
Domains

Heuristic Search

Local search

Heuristics

Properties of greedy best-first search

- one of the three most commonly used algorithms for satisficing planning
- **complete** for safe heuristics (due to duplicate detection)
- **suboptimal** unless h satisfies some very strong assumptions (similar to being perfect)
- invariant under all strictly monotonic transformations of h (e.g., scaling with a positive constant or adding a constant)

Reinforcement
Learning

(SDMRL)

Sarah Keren

Model-free vs.
Model-based
Decision-making

Recipes for
Control
Approaches

Decision-making
as Supervised
Learning

Model-based
Planning

Planning for
Deterministic
Domains

Heuristic Search

Local search

Heuristics

Chooses next node to expand based on $f(n) = g(n) + h(n)$

- 1 $g(n)$ Distance from start
- 2 $h(n)$ Heuristic function that estimated the expected distance from goal

A heuristic is *admissible* if it 'optimisitic': it underestimates the cost to goal.

Key points:

- As long as the heuristic is admissible then A* an optimal solution.
- Trade-of between estimation quality and computation cost.
- h = straight-line / Manhattan distance is a good heuristic for motion planning.

A* (with duplicate detection and reopening)

```

open := new min-heap ordered by  $(\sigma \mapsto f(\sigma) = g(\sigma) + h(\sigma))$ 
open.insert(make-root-node(init()))
closed :=  $\emptyset$ 
distance :=  $\emptyset$ 
while not open.empty():
     $\sigma = \text{open.pop-min}()$ 
    if  $\text{state}(\sigma) \notin \text{closed}$  or  $g(\sigma) < \text{distance}(\text{state}(\sigma))$ :
        closed := closed  $\cup \{\text{state}(\sigma)\}$ 
        distance(state( $\sigma$ )) :=  $g(\sigma)$ 
        if is-goal(state( $\sigma$ )):
            return extract-solution( $\sigma$ )
        for each  $\langle o, s \rangle \in \text{succ}(\text{state}(\sigma))$ :
             $\sigma' := \text{make-node}(\sigma, o, s)$ 
            if  $h(\sigma') < \infty$ :
                open.insert( $\sigma'$ )
return unsolvable

```

A*(one node maintained per state)

A* (with duplicate detection and reopening)

```
open := new min-set-heap ordered by  $(\sigma \mapsto f(\sigma) = g(\sigma) + h(\sigma))$ 
open.insert(make-root-node(init()))
closed :=  $\emptyset$ 
distance :=  $\emptyset$ 
while not open.empty():
     $\sigma = \text{open.pop-min}()$ 
    if  $\text{state}(\sigma) \notin \text{closed}$  or  $g(\sigma) < \text{distance}(\text{state}(\sigma))$ :
        closed := closed  $\cup$  {state( $\sigma$ )}
        distance(state( $\sigma$ )) :=  $g(\sigma)$ 
        if is-goal(state( $\sigma$ )):
            return extract-solution( $\sigma$ )
        for each  $\langle o, s \rangle \in \text{succ}(\text{state}(\sigma))$ :
             $\sigma' := \text{make-node}(\sigma, o, s)$ 
            if  $h(\sigma') < \infty$ :
                open.insert( $\sigma'$ )
return unsolvable
```

Reinforcement
Learning

(SDMRL)

Sarah Keren

Model-free vs.
Model-based
Decision-making

Recipes for
Control
Approaches

Decision-making
as Supervised
Learning

Model-based
Planning

Planning for
Deterministic
Domains

Heuristic Search

Local search

Heuristics

A* by example


(a) The initial state




(b) After expanding Arad



(c) After expanding Sibiu



(d) After expanding Rimnicu Vilcea



(e) After expanding Fagaras

Reinforcement Learning

(SDMRL)

Sarah Keren

Model-free vs. Model-based Decision-making

Recipes for Control Approaches

Decision-making as Supervised Learning

Model-based Planning

Planning for Deterministic Domains


Heuristic Search

Local search

Heuristics

Sarah Keren

A* by example



Reinforcement
Learning

(SDMRL)

Sarah Keren

Model-free vs.
Model-based
Decision-making

Recipes for
Control
Approaches

Decision-making
as Supervised
Learning

Model-based
Planning

Planning for
Deterministic
Domains

Heuristic Search

Local search

Heuristics

Properties of A*

- the most commonly used algorithm for optimal planning
- rarely used for satisficing planning
- **complete** for safe heuristics
(even without duplicate detection)
- **optimal** for admissible heuristics

Reinforcement
Learning

(SDMRL)

Sarah Keren

Model-free vs.
Model-based
Decision-making

Recipes for
Control
Approaches

Decision-making
as Supervised
Learning

Model-based
Planning


Planning for
Deterministic
Domains

Heuristic Search

Local search

Heuristics

Optimality of Tree-Search A* with admissible h



Reinforcement
Learning

(SDMRL)

Sarah Keren

Model-free vs.
Model-based
Decision-making

Recipes for
Control
Approaches

Decision-making
as Supervised
Learning

Model-based
Planning

Planning for
Deterministic
Domains

Heuristic Search

Local search

Heuristics

Optimality of Tree-Search A* with admissible h

Suppose to the contrary that the algorithm returns a suboptimal plan via *extract-solution*(σ) for some node σ with $state(\sigma) \in G$.

- If so, then no optimal goal node σ^* was in the open list (right?)
- If so, then open list contains some unexpanded node σ' on the (optimal) path from root node to σ^*

Reinforcement
Learning

(SDMRL)

Sarah Keren

Model-free vs.
Model-based
Decision-making

Recipes for
Control
Approaches

Decision-making
as Supervised
Learning

Model-based
Planning

Planning for
Deterministic
Domains

Heuristic Search

Local search

Heuristics

Optimality of Tree-Search A* with admissible h

Suppose to the contrary that the algorithm returns a suboptimal plan via *extract-solution*(σ) for some node σ with $state(\sigma) \in G$.

- If so, then no optimal goal node σ^* was in the open list (right?)
- If so, then open list contains some unexpanded node σ' on the (optimal) path from root node to σ^*

$$\begin{aligned} f(\sigma) &= g(\sigma) && \text{since } h(\sigma) = 0 \\ &> g(\sigma^*) && \text{since } \sigma \text{ is suboptimal} \\ &\geq f(\sigma') && \text{since } h \text{ is admissible} \end{aligned}$$

$f(\sigma) > f(\sigma') \rightsquigarrow$ contradiction with “ σ is dequeued before σ' ”

Weighted A*

Weighted A* (with duplicate detection and reopening)

$open := \text{new min-heap ordered by } (\sigma \mapsto g(\sigma) + W \cdot h(\sigma))$

$open.insert(\text{make-root-node}(\text{init()}))$

$closed := \emptyset$

$distance := \emptyset$

while not $open.empty()$:

$\sigma = open.pop\text{-}min()$

if $state(\sigma) \notin closed$ **or** $g(\sigma) < distance(state(\sigma))$:

$closed := closed \cup \{state(\sigma)\}$

$distance(\sigma) := g(\sigma)$

if $\text{is-goal}(state(\sigma))$:

return $\text{extract-solution}(\sigma)$

for each $\langle o, s \rangle \in \text{succ}(state(\sigma))$:

$\sigma' := \text{make-node}(\sigma, o, s)$

if $h(\sigma') < \infty$:

$open.insert(\sigma')$

return unsolvable

Reinforcement
Learning

(SDMRL)

Sarah Keren

Model-free vs.
Model-based
Decision-making

Recipes for
Control
Approaches

Decision-making
as Supervised
Learning

Model-based
Planning

Planning for
Deterministic
Domains

Heuristic Search

Local search

Heuristics

Properties of weighted A*

The **weight** $W \in \mathbb{R}_0^+$ is a parameter of the algorithm.

- for $W = 0$, behaves like breadth-first search
- for $W = 1$, behaves like A*
- for $W \rightarrow \infty$, behaves like greedy best-first search

Properties:

- one of the three most commonly used algorithms for satisficing planning
- for $W > 1$, can prove similar properties to A*, replacing **optimal** with **bounded suboptimal**: generated solutions are at most a factor W as long as optimal ones

Reinforcement
Learning

(SDMRL)

Sarah Keren

Model-free vs.
Model-based
Decision-making

Recipes for
Control
Approaches

Decision-making
as Supervised
Learning

Model-based
Planning

Planning for
Deterministic
Domains

Heuristic Search

Local search



Heuristics

Algorithm 1 Best First Search (BFS)

BFS(P, h)

```
1: create OPEN list for unexpanded nodes
2: put  $\langle P.root, h(P.root) \rangle$  in OPEN
3:  $n_{cur} = ExtractMax(OPEN)$  (initial model)
4: while  $n_{cur}$  do
5:   if  $IsTerminal(n_{cur})$  then
6:     return  $ExtractPath(n_{cur})$  (best solution found - exit)
7:   end if
8:   for all  $n_{suc} \in GetSuccessors(n_{cur}, P)$  do
9:     put  $\langle n_{suc}, h(n_{suc}) \rangle$  in OPEN
10:  end for
11:   $n_{cur} = ExtractMax(OPEN)$ 
12: end while
13: return no solution
```

Best First Design



Reinforcement
Learning

(SDMRL)

Sarah Keren

Model-free vs.
Model-based
Decision-making

Recipes for
Control
Approaches

Decision-making
as Supervised
Learning

Model-based
Planning

Planning for
Deterministic
Domains

Heuristic Search

Local search

Heuristics


Algorithm 2 Best First Design (BFD)

BFD(δ, h)

```
1: create OPEN list for unexpanded nodes
2:  $n_{cur} = \langle design, \vec{m}_{\emptyset} \rangle$  (initial model)
3: while  $n_{cur}$  do
4:   if  $IsExecution(n_{cur})$  then
5:     return  $n_{cur}.\vec{m}$  (best modification found - exit)
6:   end if
7:   for all  $n_{suc} \in GetSuccessors(n_{cur}, \delta)$  do
8:     put  $\langle \langle design, n_{suc}.\vec{m} \rangle, h(n_{suc}) \rangle$  in OPEN
9:   end for
10:  if  $\Phi_{\sigma}(n_{cur}.\vec{m}) = 1$  then
11:    put  $\langle \langle execution, \vec{m}_{new} \rangle, v^*(\delta_{\vec{m}_{new}}) \rangle$  in OPEN
12:  end if
13:   $n_{cur} = ExtractMax(OPEN)$ 
14: end while
15: return error
```

Hill-climbing

A local search algorithm that incrementally improves the solution by exploring the neighboring states of the current state.



Reinforcement
Learning

(SDMRL)

Sarah Keren

Model-free vs.
Model-based
Decision-making

Recipes for
Control
Approaches

Decision-making
as Supervised
Learning

Model-based
Planning

Planning for
Deterministic
Domains

Heuristic Search

Local search

Heuristics

Hill-climbing

```
 $\sigma := \text{make-root-node}(\text{init}())$   
forever:  
  if is-goal(state( $\sigma$ )):  
    return extract-solution( $\sigma$ )  
   $\Sigma' := \{ \text{make-node}(\sigma, o, s) \mid \langle o, s \rangle \in \text{succ}(\text{state}(\sigma)) \}$   
   $\sigma :=$  an element of  $\Sigma'$  minimizing  $h$  (random tie breaking)
```

- can get stuck in **local minima** where immediate improvements of $h(\sigma)$ are not possible
- many variations: tie-breaking strategies, restarts

Enforced hill-climbing

Enforced hill-climbing: procedure improve

```
def improve( $\sigma_0$ ):  
    queue := new fifo-queue  
    queue.push-back( $\sigma_0$ )  
    closed :=  $\emptyset$   
    while not queue.empty():  
         $\sigma$  = queue.pop-front()  
        if state( $\sigma$ )  $\notin$  closed:  
            closed := closed  $\cup$  {state( $\sigma$ )}  
            if  $h(\sigma) < h(\sigma_0)$ :  
                return  $\sigma$   
            for each  $\langle o, s \rangle \in \text{succ}(\text{state}(\sigma))$ :  
                 $\sigma' := \text{make-node}(\sigma, o, s)$   
                queue.push-back( $\sigma'$ )  
    fail
```

\leadsto breadth-first search for more promising node than σ_0

Reinforcement
Learning

(SDMRL)

Sarah Keren

Model-free vs.
Model-based
Decision-making

Recipes for
Control
Approaches

Decision-making
as Supervised
Learning

Model-based
Planning

Planning for
Deterministic
Domains

Heuristic Search

Local search

Heuristics

Enforced hill-climbing (ctd.)

Enforced hill-climbing

```
 $\sigma := \text{make-root-node}(\text{init}())$   
while not is-goal(state( $\sigma$ )):  
     $\sigma := \text{improve}(\sigma)$   
return extract-solution( $\sigma$ )
```

- one of the three most commonly used algorithms for satisficing planning
- can fail if procedure *improve* fails (when the goal is unreachable from σ_0)
- complete for **undirected** search spaces (where the successor relation is symmetric) if $h(\sigma) = 0$ for all goal nodes and only for goal nodes

Classification: what works where in (deterministic) planning?

uninformed vs. heuristic search:

- For **satisficing** planning, heuristic search vastly outperforms uninformed algorithms on most domains.
- For **optimal** planning, heuristic search typically outperforms uninformed algorithms, but an efficiently implemented uninformed algorithm is not easy to beat in most domains.

systematic search vs. local search:

- For **satisficing** planning, the most successful algorithms are somewhere between the two extremes.
- For **optimal** planning, systematic algorithms are required.

Reinforcement
Learning

(SDMRL)

Sarah Keren

Model-free vs.
Model-based
Decision-making

Recipes for
Control
Approaches

Decision-making
as Supervised
Learning

Model-based
Planning

Planning for
Deterministic
Domains

Heuristic Search

Local search

Heuristics

Where heuristics come from?

General idea

(Admissible) heuristic functions obtained as
(optimal) cost functions of relaxed problems

Examples

- Euclidian distance in Path Finding
- Manhattan distance in N-puzzle
- Spanning Tree in Traveling Salesman Problem
- Shortest Path in Job Shop Scheduling

Reinforcement
Learning

(SDMRL)

Sarah Keren

Model-free vs.
Model-based
Decision-making

Recipes for
Control
Approaches

Decision-making
as Supervised
Learning

Model-based
Planning

Planning for
Deterministic
Domains


Heuristic Search

Local search


Heuristics

Example

True distance h^* for different search states



Manhattan distance is based on the relaxation that ...?



Example

Reinforcement
Learning

(SDMRL)

Sarah Keren

Model-free vs.
Model-based
Decision-making

Recipes for
Control
Approaches

Decision-making
as Supervised
Learning

Model-based
Planning

Planning for
Deterministic
Domains

Heuristic Search

Local search

Heuristics

Example

7	2	4
5		6
8	3	1


Start State

1	2	3
4	5	6
7	8	

Goal State

- A tile can move from square A to square B if A is adjacent to B and B is blank \leadsto solution distance h^*
- A tile can move from square A to square B if A is adjacent to B \leadsto manhattan distance heuristic h^{MD}
- A tile can move from square A to square B \leadsto misplaced tiles heuristic h^{MT}

Example



- A tile can move from square A to square B if A is adjacent to B and B is blank \leadsto solution distance h^*
- A tile can move from square A to square B if A is adjacent to B \leadsto manhattan distance heuristic h^{MD}
- A tile can move from square A to square B \leadsto misplaced tiles heuristic h^{MT}

Here: $h^*(s_0) = ?$, $h^{MD}(s_0) = 14$, $h^{MT}(s_0) = 6$

In general, $h^* \geq h^{MD} \geq h^{MT}$. (Why?)

Relaxations can be used for heuristic estimations.

Different possibilities:

- Implement an **optimal planner** for relaxed planning tasks and use its solution lengths as an estimate, even though it is NP-hard.
 \leadsto **h^+ heuristic** -ignore delete effects.
- Do not actually solve the relaxed planning task, but compute an estimate of its difficulty in a different way.
 \leadsto **h_{\max} heuristic, h_{add} heuristic**
- Compute a solution for relaxed planning tasks which is not necessarily optimal, but “reasonable”.
 \leadsto **h_{FF} heuristic**


Graphical “interpretation”: Relaxed planning graphs

- Build a layered **reachability graph** $P_0, A_0, P_1, A_1, \dots$

$$P_0 = \{p \in I\}$$

$$A_i = \{a \in A \mid \text{pre}(a) \subseteq P_i\}$$

$$P_{i+1} = P_i \cup \{p \in \text{add}(a) \mid a \in A_i\}$$



- Terminate when $G \subseteq P_i$

Running example

$$P = \{a, b, c, d, e, f, g, h\}$$

$$I = \{a\}$$

$$G = \{c, d, e, f, g\}$$

$$a_1 = \langle \{a\}, \{b, c\}, \{a\} \rangle$$

$$a_2 = \langle \{a, c\}, \{d\}, \{d\} \rangle$$

$$a_3 = \langle \{b, c\}, \{e\}, \{e, f\} \rangle$$

$$a_4 = \langle \{b\}, \{f\}, \emptyset \rangle$$

$$a_5 = \langle \{d\}, \{e, f\}, \{d\} \rangle$$

$$a_6 = \langle \{d\}, \{g\}, \{b\} \rangle$$

Running example

$$P = \{a, b, c, d, e, f, g, h\}$$

$$I = \{a\}$$

$$G = \{c, d, e, f, g\}$$

$$a_1 = \langle \{a\}, \{b, c\}, \emptyset \rangle$$

$$a_2 = \langle \{a, c\}, \{d\}, \emptyset \rangle$$


$$a_3 = \langle \{b, c\}, \{e\}, \emptyset \rangle$$

$$a_4 = \langle \{b\}, \{f\}, \emptyset \rangle$$

$$a_5 = \langle \{d\}, \{e, f\}, \emptyset \rangle$$

$$a_6 = \langle \{d\}, \{g\}, \emptyset \rangle$$

Running example: Relaxed planning graph



Reinforcement
Learning

(SDMRL)

Sarah Keren

Model-free vs.
Model-based
Decision-making

Recipes for
Control
Approaches

Decision-making
as Supervised
Learning

Model-based
Planning


Planning for
Deterministic
Domains

Heuristic Search

Local search

Heuristics

Running example: Relaxed planning graph



Reinforcement
Learning

(SDMRL)

Sarah Keren

Model-free vs.
Model-based
Decision-making

Recipes for
Control
Approaches

Decision-making
as Supervised
Learning

Model-based
Planning

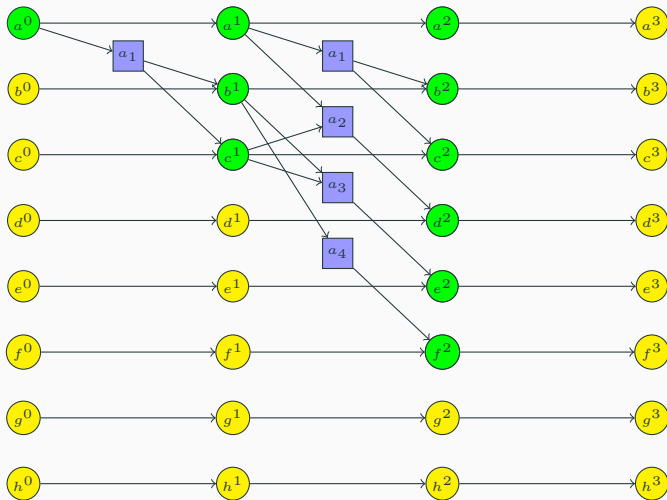
Planning for
Deterministic
Domains

Heuristic Search

Local search

Heuristics

Running example: Relaxed planning graph



Reinforcement
Learning

(SDMRL)

Sarah Keren

Model-free vs.
Model-based
Decision-making

Recipes for
Control
Approaches

Decision-making
as Supervised
Learning

Model-based
Planning

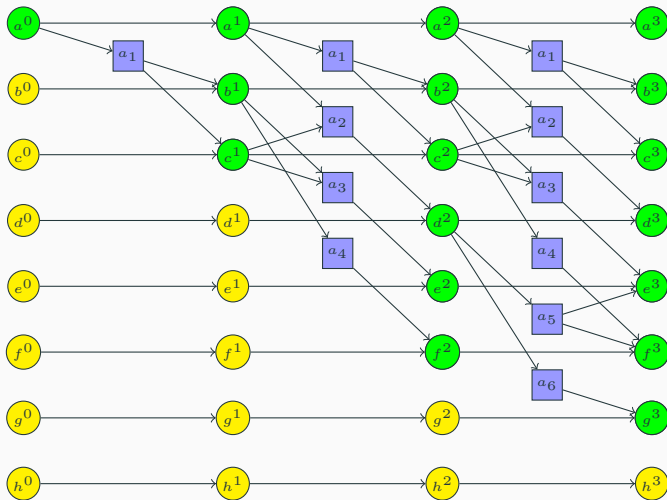
Planning for
Deterministic
Domains

Heuristic Search

Local search

Heuristics

Running example: Relaxed planning graph



Reinforcement
Learning

(SDMRL)

Sarah Keren

Model-free vs.
Model-based
Decision-making

Recipes for
Control
Approaches

Decision-making
as Supervised
Learning

Model-based
Planning

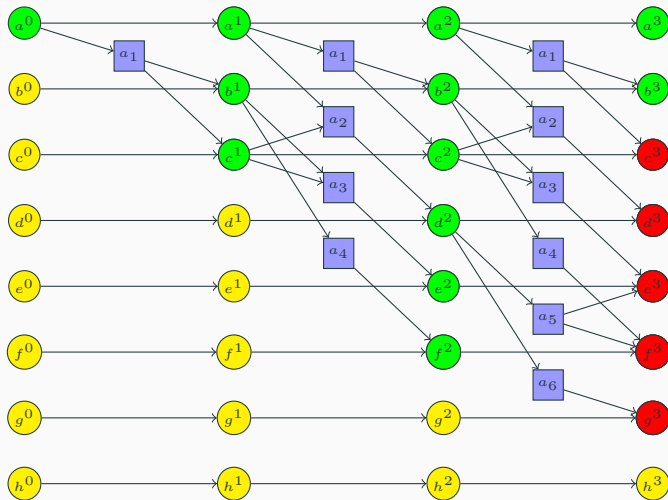
Planning for
Deterministic
Domains

Heuristic Search

Local search

Heuristics

Running example: Relaxed planning graph



Reinforcement
Learning

(SDMRL)

Sarah Keren

Model-free vs.
Model-based
Decision-making

Recipes for
Control
Approaches

Decision-making
as Supervised
Learning

Model-based
Planning

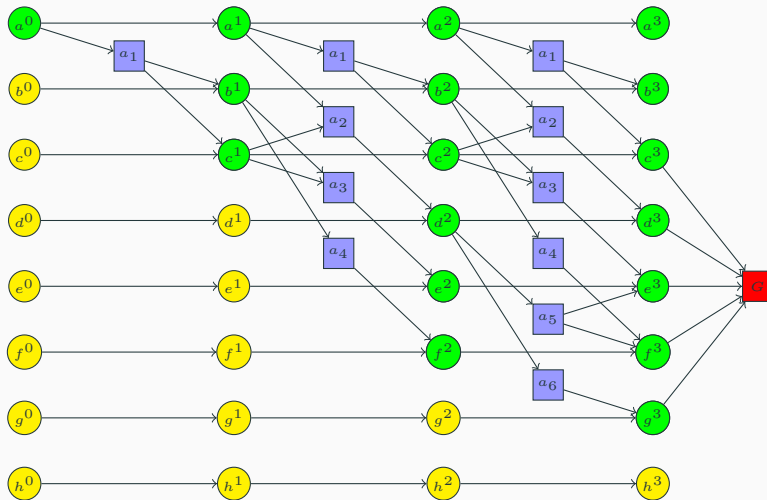
Planning for
Deterministic
Domains

Heuristic Search

Local search

Heuristics

Running example: Relaxed planning graph



Reinforcement
Learning

(SDMRL)

Sarah Keren

Model-free vs.
Model-based
Decision-making

Recipes for
Control
Approaches

Decision-making
as Supervised
Learning

Model-based
Planning

Planning for
Deterministic
Domains

Heuristic Search

Local search

Heuristics

Dominance relation between admissible heuristics

Precision matters

Given two admissible heuristics h_1, h_2 , if $h_2(\sigma) \geq h_1(\sigma)$ for all search nodes σ , then h_2 **dominates** h_1 and is better for optimizing search

Typical search costs (unit-cost action)

$h^*(I) = 14$	BFS $\approx 1,700,000$ nodes
	$A^*(h_1) \approx 560$ nodes
	$A^*(h_2) \approx 115$ nodes
$h^*(I) = 24$	BFS $\approx 27,000,000,000$ nodes
	$A^*(h_1) \approx 40,000$ nodes
	$A^*(h_2) \approx 1,650$ nodes

Reinforcement
Learning

(SDMRL)

Sarah Keren

Model-free vs.
Model-based
Decision-making

Recipes for
Control
Approaches

Decision-making
as Supervised
Learning

Model-based
Planning

Planning for
Deterministic
Domains

Heuristic Search
Local search

Heuristics

Dominance relation between admissible heuristics

Precision matters

Given two admissible heuristics h_1, h_2 , if $h_2(\sigma) \geq h_1(\sigma)$ for all search nodes σ , then h_2 **dominates** h_1 and is better for optimizing search

Combining admissible heuristics

For any admissible heuristics h_1, \dots, h_k ,

$$h(\sigma) = \max_{i=1}^k \{h_i(\sigma)\}$$

is also admissible and dominates all individual h_i

Reinforcement
Learning

(SDMRL)

Sarah Keren

Model-free vs.
Model-based
Decision-making

Recipes for
Control
Approaches

Decision-making
as Supervised
Learning

Model-based
Planning

Planning for
Deterministic
Domains

Heuristic Search

Local search

Heuristics

Are we done?

General idea

(Admissible) heuristic functions obtained as (optimal) cost functions of relaxed problems

- OK, but heuristic is **yet another input** to our agent!
- Satisfactory for general solvers?
- Satisfactory in special purpose solvers?

Reinforcement
Learning

(SDMRL)

Sarah Keren

Model-free vs.
Model-based
Decision-making

Recipes for
Control
Approaches

Decision-making
as Supervised
Learning

Model-based
Planning

Planning for
Deterministic
Domains

Heuristic Search

Local search

Heuristics

Are we done?

General idea

(Admissible) heuristic functions obtained as (optimal) cost functions of relaxed problems

- OK, but heuristic is **yet another input** to our agent!
- Satisfactory for general solvers?
- Satisfactory in special purpose solvers?

Towards domain-independent agents

- How to get heuristics **automatically**?
- Can such automatically derived heuristics **dominate** the domain-specific heuristics crafted by hand?

Reinforcement
Learning

(SDMRL)

Sarah Keren

Model-free vs.
Model-based
Decision-making

Recipes for
Control
Approaches

Decision-making
as Supervised
Learning

Model-based
Planning

Planning for
Deterministic
Domains

Heuristic Search

Local search

Heuristics

Example Heuristics



Reinforcement
Learning

(SDMRL)

Sarah Keren

Model-free vs.
Model-based
Decision-making

Recipes for
Control
Approaches

Decision-making
as Supervised
Learning

Model-based
Planning

Planning for
Deterministic
Domains

Heuristic Search

Local search

Heuristics

Recap

- Spectrum of approaches to control
- Various characteristics (model-free model-based, offlineonline, etc)
- Policy extraction as supervised learning
- Model-based planning in deterministic and full observable domains



Reinforcement
Learning

(SDMRL)

Sarah Keren

Model-free vs.
Model-based
Decision-making

Recipes for
Control
Approaches

Decision-making
as Supervised
Learning

Model-based
Planning

Planning for
Deterministic
Domains

Heuristic Search

Local search

Heuristics