

Sequential Decision Making and Reinforcement Learning

(SDMRL)

Model-free RL

Sarah Keren

The Taub Faculty of Computer Science
Technion - Israel Institute of Technology

Acknowledgments

- David Sliver's course on RL:
<https://www.deepmind.com/learning-resources/introduction-to-reinforcement-learning-with-david-silver>
- Slides by Malte Helmert, Carmel Domshlak, Erez Karpas and Alexander Shleyfman.

Reinforcement
Learning

(SDMRL)

Sarah Keren

Recap

Model Free RL:
Monte Carlo

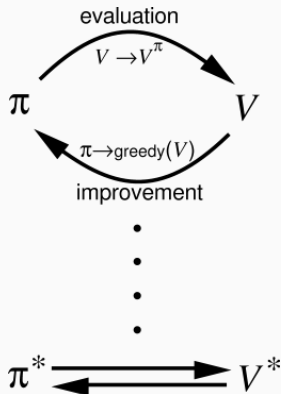
Temporal-
Difference
Learning

Comparing Monte
Carlo and TD
methods

Policy Search
Methods

Recap

Anatomy of RL Algorithms



Reinforcement
Learning

(SDMRL)

Sarah Keren

Recap

Model Free RL:
Monte Carlo

Temporal-
Difference
Learning

Comparing Monte
Carlo and TD
methods

Policy Search
Methods

Model-Free vs. Model-Based RL




Figure 1: By Michael Littman

<https://www.youtube.com/watch?v=45FKxa3qPHo>

Model-Free vs. Model-Based RL



Reinforcement
Learning

(SDMRL)

Sarah Keren

Recap

Model Free RL:
Monte Carlo

Temporal-
Difference
Learning


Comparing Monte
Carlo and TD
methods

Policy Search
Methods

Model Free RL: Monte Carlo

Monte Carlo (MC)

- MC methods learn directly from episodes of experience
- MC is model-free: no knowledge of MDP transitions / rewards
- MC learns from complete episodes: no bootstrapping.
- MC uses the simplest possible idea: value = mean return
- Caveat: can only apply MC to episodic MDPs: all episodes must terminate



Reinforcement
Learning

(SDMRL)

Sarah Keren

Recap

Model Free RL:
Monte Carlo

Temporal-
Difference
Learning

Comparing Monte
Carlo and TD
methods

Policy Search
Methods

Monte Carlo Simulation-Rollout

- A Monte Carlo simulation is a statistical technique used to analyze the behavior of complex systems and processes that are influenced by random variables.

Reinforcement
Learning

(SDMRL)

Sarah Keren

Recap

Model Free RL:
Monte Carlo

Temporal-
Difference
Learning

Comparing Monte
Carlo and TD
methods

Policy Search
Methods

Monte Carlo Simulation-Rollout

- A Monte Carlo simulation is a statistical technique used to analyze the behavior of complex systems and processes that are influenced by random variables.
- Each rollout involves simulating a sequence of actions, typically by selecting actions according to some policy (e.g., uniformly at random, or based on a heuristic) until a terminal state is reached.

Reinforcement
Learning

(SDMRL)

Sarah Keren

Recap

Model Free RL:
Monte Carlo

Temporal-
Difference
Learning

Comparing Monte
Carlo and TD
methods

Policy Search
Methods

Monte Carlo Simulation-Rollout

- A Monte Carlo simulation is a statistical technique used to analyze the behavior of complex systems and processes that are influenced by random variables.
- Each rollout involves simulating a sequence of actions, typically by selecting actions according to some policy (e.g., uniformly at random, or based on a heuristic) until a terminal state is reached.
- Typically, a large number of simulations are used to estimate the probabilities and distributions of different outcomes.

Reinforcement
Learning

(SDMRL)

Sarah Keren

Recap

Model Free RL:
Monte Carlo

Temporal-
Difference
Learning

Comparing Monte
Carlo and TD
methods

Policy Search
Methods

Monte Carlo Simulation-Rollout

- A Monte Carlo simulation is a statistical technique used to analyze the behavior of complex systems and processes that are influenced by random variables.
- Each rollout involves simulating a sequence of actions, typically by selecting actions according to some policy (e.g., uniformly at random, or based on a heuristic) until a terminal state is reached.
- Typically, a large number of simulations are used to estimate the probabilities and distributions of different outcomes.
- We will use this in different algorithms



Reinforcement
Learning

(SDMRL)

Sarah Keren

Recap

Model Free RL:
Monte Carlo

Temporal-
Difference
Learning

Comparing Monte
Carlo and TD
methods

Policy Search
Methods

- **Reminder:**

- Return is (typically) the total discounted reward:

$$G_t = R_{t+1} + \gamma R_{t+2} + \dots + \gamma^{T-1} R_T$$

- Value function is (typically) the expected return:

$$V_{\pi}(s) = \mathbb{E}_{\pi}[G_t | S_t = s]$$

- **Objective:** learn V_{π} from episodes of experience under policy π

$$S_1, A_1, R_2, \dots, S_k \sim \pi$$

- Estimate $V_{\pi}(s)$ as average total reward of epochs containing s (calculating from s to end of epoch).

Monte-Carlo Policy Evaluation

- **Reminder:**

- Return is (typically) the total discounted reward:

$$G_t = R_{t+1} + \gamma R_{t+2} + \dots + \gamma^{T-1} R_T$$

- Value function is (typically) the expected return:

$$V_\pi(s) = \mathbb{E}_\pi[G_t | S_t = s]$$

- **Objective:** learn V_π from episodes of experience under policy π

$$S_1, A_1, R_2, \dots, S_k \sim \pi$$

- Estimate $V_\pi(s)$ as average total reward of epochs containing s (calculating from s to end of epoch).

Monte-Carlo policy evaluation uses empirical mean return instead of expected return. **Why ?**

Monte-Carlo Policy Evaluation

Key Idea

Use observed reward-to-go of the state as the direct evidence of the actual expected utility of that state.

- Reward-to-go of a state s = the sum of the (discounted) rewards from that state until a terminal state is reached
- Two versions to evaluate state s :
 - The first time-step t that state s is visited in an episode
 - Every time-step t that state s is visited in an episode
- Increment visit counter $N(s) \leftarrow N(s) + 1$
- Increment total return $S(s) \leftarrow S(s) + G_t$
- Value is estimated by mean return $V(s) = S(s)/N(s)$
- By law of large numbers, $V(s) \rightarrow V_\pi(s)$ as $N(s) \rightarrow \infty$ (averaging the reward-to-go samples will converge to true value at state)

Reinforcement
Learning

(SDMRL)

Sarah Keren

Recap

Model Free RL:
Monte Carlo

Temporal-
Difference
Learning

Comparing Monte
Carlo and TD
methods

Policy Search
Methods

Monte-Carlo Policy Evaluation - Blackjack

- States (200 of them):
 - Current sum (12-21)
 - Dealer's showing card (ace-10)
 - Do I have a "useable" ace? (yes-no)
- Action stick: Stop receiving cards (and terminate)
- Action twist: Take another card (no replacement)
- Reward for stick:
 - +1 if sum of cards > sum of dealer cards
 - 0 if sum of cards = sum of dealer cards
 - -1 if sum of cards < sum of dealer cards
- Reward for twist:
 - -1 if sum of cards > 21 (and terminate)
 - 0 otherwise
- Transitions: automatically twist if sum of cards < 12



Reinforcement
Learning

(SDMRL)

Sarah Keren

Recap

Model Free RL:
Monte Carlo

Temporal-
Difference
Learning

Comparing Monte
Carlo and TD
methods

Policy Search
Methods

Monte-Carlo Policy Evaluation

Reinforcement
Learning

(SDMRL)

Sarah Keren

Recap


Model Free RL:
Monte Carlo

Temporal-
Difference
Learning

Comparing Monte
Carlo and TD
methods

Policy Search
Methods

Blackjack Value Function after Monte-Carlo Learning



Monte-Carlo Policy Evaluation

Reinforcement
Learning

(SDMRL)

Sarah Keren

Recap

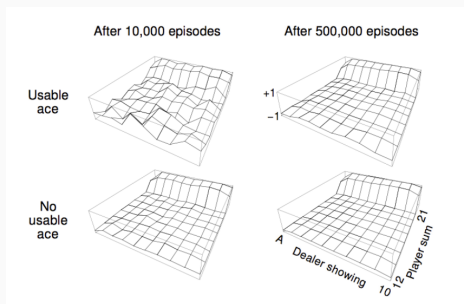
Model Free RL:
Monte Carlo

Temporal-
Difference
Learning

Comparing Monte
Carlo and TD
methods

Policy Search
Methods

Blackjack Value Function after Monte-Carlo Learning



Policy - stick if sum of cards ≥ 20 , otherwise twist.

Prediction:

Initialize:

$\pi \leftarrow$ policy to be evaluated

$V \leftarrow$ an arbitrary state-value function

$Returns(s) \leftarrow$ an empty list, for all $s \in \mathcal{S}$

Repeat forever:

Generate an episode using π

For each state s appearing in the episode:

$G \leftarrow$ return following the first occurrence of s

Append G to $Returns(s)$

$V(s) \leftarrow \text{average}(Returns(s))$

Prediction:

Initialize:

$\pi \leftarrow$ policy to be evaluated

$V \leftarrow$ an arbitrary state-value function

$Returns(s) \leftarrow$ an empty list, for all $s \in \mathcal{S}$

Repeat forever:

Generate an episode using π

For each state s appearing in the episode:

$G \leftarrow$ return following the first occurrence of s

Append G to $Returns(s)$

$V(s) \leftarrow \text{average}(Returns(s))$

How can we use this for control ?

Control:

Initialize, for all $s \in \mathcal{S}$, $a \in \mathcal{A}(s)$:

$Q(s, a) \leftarrow \text{arbitrary}$

$\pi(s) \leftarrow \text{arbitrary}$

$Returns(s, a) \leftarrow \text{empty list}$

Repeat forever:

Choose $S_0 \in \mathcal{S}$ and $A_0 \in \mathcal{A}(S_0)$ s.t. all pairs have probability > 0

Generate an episode starting from S_0, A_0 , following π

For each pair s, a appearing in the episode:

$G \leftarrow \text{return following the first occurrence of } s, a$

Append G to $Returns(s, a)$

$Q(s, a) \leftarrow \text{average}(Returns(s, a))$

For each s in the episode:

$\pi(s) \leftarrow \arg\max_a Q(s, a)$

Control:

Initialize, for all $s \in \mathcal{S}$, $a \in \mathcal{A}(s)$:

$Q(s, a) \leftarrow \text{arbitrary}$

$\pi(s) \leftarrow \text{arbitrary}$

$Returns(s, a) \leftarrow \text{empty list}$

Repeat forever:

Choose $S_0 \in \mathcal{S}$ and $A_0 \in \mathcal{A}(S_0)$ s.t. all pairs have probability > 0

Generate an episode starting from S_0, A_0 , following π

For each pair s, a appearing in the episode:

$G \leftarrow \text{return following the first occurrence of } s, a$

Append G to $Returns(s, a)$

$Q(s, a) \leftarrow \text{average}(Returns(s, a))$

For each s in the episode:

$\pi(s) \leftarrow \arg\max_a Q(s, a)$

Problem with this approach?

Monte-Carlo: From Prediction (Evaluation) to Control

Reinforcement
Learning

(SDMRL)

Sarah Keren

Recap

Model Free RL:
Monte Carlo

Temporal-
Difference
Learning

Comparing Monte
Carlo and TD
methods

Policy Search
Methods

- Use a random policy to simulate many (!) trajectories
- Compute the q-value of each state-action pair
- Update π by taking the max action.

Problem with this approach?

Monte-Carlo Policy Control - Taxi

```
+-----+
|R: | : :G| |
| : | : : |
| : | : : |
| | : | : |
|Y| : |B: |
+-----+
(North)
num episodes completed: 1
total rewards: -133
mean rewards per episode: -133.00
```

```
def build_decision_dict(raw_data):
    # Nested dictionary for: State -> Action -> Reward List
    state_action_scores = defaultdict(lambda: defaultdict(lambda: []))
    for trajectory in raw_data:
        reward_sum = 0
        # iterate backwards to calculate the return G of each observed state action pair
        for state, action, reward in reversed(list(zip(trajectory.observations, trajectory.actions, trajectory.rewards))):
            reward_sum += reward
            state_action_scores[state][action].append(reward_sum)

    for state, action_values in state_action_scores.items():
        for action, values_list in action_values.items():
            # Calculate the mean of all returns for a state action pair
            state_action_scores[state][action] = np.mean(values_list)
        # For each state choose the action with the highest mean return
        state_action_scores[state] = max(state_action_scores[state], key=state_action_scores[state].get)
    return state_action_scores
```

https://github.com/CLAIR-LAB-TECHNION/FSTMA-course/blob/main/tutorials/tut03/Monte_Carlo.ipynb

Reinforcement
Learning

(SDMRL)

Sarah Keren

Recap

Model Free RL:
Monte Carlo

Temporal-
Difference
Learning

Comparing Monte
Carlo and TD
methods

Policy Search
Methods

Monte-Carlo Control - Taxi

```
env.seed(seed)
```

```
policy = calc_final_policy(RandomTraveler, 100, "mcc_100")
evaluate_and_print(policy)
policy = calc_final_policy(RandomTraveler, 1000, "mcc_1000")
evaluate_and_print(policy)
policy = calc_final_policy(RandomTraveler, 10000, "mcc_10000")
evaluate_and_print(policy)
policy = calc_final_policy(RandomTraveler, 100000, "mcc_100000")
evaluate_and_print(policy)
```

```
100% ██████████ 10000/10000 [00:01<00:00, 5139.08it/s]
Monte Carlo Control Policy
-----
total reward over all episodes: -1306708
mean reward per episode:      -130.6708

100% ██████████ 10000/10000 [00:01<00:00, 5334.19it/s]
Monte Carlo Control Policy
-----
total reward over all episodes: -1242335
mean reward per episode:      -124.2335

100% ██████████ 10000/10000 [00:01<00:00, 6415.93it/s]
Monte Carlo Control Policy
-----
total reward over all episodes: -729155
mean reward per episode:      -72.9155

100% ██████████ 10000/10000 [00:01<00:00, 7840.29it/s]
Monte Carlo Control Policy
-----
total reward over all episodes: -154428
mean reward per episode:      -15.4428
```

Reinforcement
Learning

(SDMRL)

Sarah Keren

Recap


Model Free RL:
Monte Carlo

Temporal-
Difference
Learning

Comparing Monte
Carlo and TD
methods

Policy Search
Methods

Generalised Policy Iteration With Monte-Carlo Evaluation



Policy evaluation: Monte-Carlo policy evaluation, $V = \mathcal{V}_\pi$?

Policy improvement: Greedy policy improvement?

Model-Free Policy Iteration Using Action-Value Function

Reinforcement
Learning

(SDMRL)

Sarah Keren

Recap

Model Free RL:
Monte Carlo

Temporal-
Difference
Learning

Comparing Monte
Carlo and TD
methods

Policy Search
Methods

$$\pi'(s) = \arg \max_{a \in A} \sum_{s' \in \mathcal{S}} P(s'|s, a) [R(s, a, s') + \gamma V^\pi(s')],$$

$$\pi'(s) = \arg \max_{a \in A} Q(s, a)$$

Model-Free Policy Iteration Using Action-Value Function

Reinforcement
Learning

(SDMRL)

Sarah Keren

Recap

Model Free RL:
Monte Carlo

Temporal-
Difference
Learning

Comparing Monte
Carlo and TD
methods

Policy Search
Methods


Greedy policy improvement over $V(s)$ requires model of MDP

$$\pi'(s) = \arg \max_{a \in A} \sum_{s' \in \mathcal{S}} P(s'|s, a) [R(s, a, s') + \gamma V^\pi(s')],$$

Greedy policy improvement over $Q(s, a)$ is model-free

$$\pi'(s) = \arg \max_{a \in A} Q(s, a)$$

Generalised Policy Iteration With Monte-Carlo Evaluation



- **Policy evaluation:** Monte-Carlo policy evaluation. $Q_\pi(s, a)$
- **Policy improvement:** Greedy policy improvement. How?

Example of Greedy Action Selection

There are two doors in front of you:

You open the left door and get reward 0

$$V(\text{left}) = 0$$

You open the right door and get reward +1

$$V(\text{right}) = +1$$

You open the right door and get reward +2

$$V(\text{right}) = +2$$

You open the right door and get reward +2

$$V(\text{right}) = +2$$

.
. .
.

Are you sure you've chosen the best door?



- Simplest idea for ensuring continual exploration
- All m actions are tried with non-zero probability
- With probability $1 - \epsilon$ choose the greedy action
- With probability ϵ choose an action at random

Policy Improvement Theorem*

Reinforcement
Learning

(SDMRL)

Sarah Keren

Recap

Model Free RL:
Monte Carlo

Temporal-
Difference
Learning

Comparing Monte
Carlo and TD
methods

Policy Search
Methods

Theorem

Let π and π' be any pair of deterministic policies. If for all $s \in \mathcal{S}$ $Q_{\pi}(s, \pi'(s)) \geq V_{\pi}(s)$ then for all $s \in \mathcal{S}$, $V_{\pi'}(s) \geq V_{\pi}(s)$. i.e., π' is an improvement over π .

π' , that can either exploit or explore, must be at least as good as π

Theorem

For any ϵ -greedy policy π , if the ϵ -greedy policy π' with respect to Q_π is an improvement, then $V_{\pi'}(s) \geq V_\pi(s)$

$$Q_\pi(s, \pi'(s)) = \sum_{a \in \mathcal{A}} \pi'(a|s) Q_\pi(s, a)$$

ϵ -Greedy Policy Improvement*

Theorem

For any ϵ -greedy policy π , if the ϵ -greedy policy π' with respect to Q_π is an improvement, then $V_{\pi'}(s) \geq V_\pi(s)$

$$\begin{aligned} Q_\pi(s, \pi'(s)) &= \sum_{a \in \mathcal{A}} \pi'(a|s) Q_\pi(s, a) \\ &= \frac{\epsilon}{m} \sum_{a \in \mathcal{A}} Q_\pi(s, a) + (1 - \epsilon) \max_{a \in \mathcal{A}} Q_\pi(s, a) \end{aligned}$$

Reinforcement
Learning

(SDMRL)

Sarah Keren

Recap

Model Free RL:
Monte Carlo

Temporal-
Difference
Learning

Comparing Monte
Carlo and TD
methods

Policy Search
Methods

ϵ -Greedy Policy Improvement*

Theorem

For any ϵ -greedy policy π , if the ϵ -greedy policy π' with respect to Q_π is an improvement, then $V_{\pi'}(s) \geq V_\pi(s)$

$$\begin{aligned} Q_\pi(s, \pi'(s)) &= \sum_{a \in \mathcal{A}} \pi'(a|s) Q_\pi(s, a) \\ &= \frac{\epsilon}{m} \sum_{a \in \mathcal{A}} Q_\pi(s, a) + (1 - \epsilon) \max_{a \in \mathcal{A}} Q_\pi(s, a) \\ &\geq \frac{\epsilon}{m} \sum_{a \in \mathcal{A}} Q_\pi(s, a) + (1 - \epsilon) \sum_{a \in \mathcal{A}} \frac{\pi(a|s) - \frac{\epsilon}{m}}{1 - \epsilon} Q_\pi(s, a) \end{aligned}$$

Reinforcement
Learning

(SDMRL)

Sarah Keren

Recap

Model Free RL:
Monte Carlo

Temporal-
Difference
Learning

Comparing Monte
Carlo and TD
methods

Policy Search
Methods

ϵ -Greedy Policy Improvement*

Theorem


For any ϵ -greedy policy π , if the ϵ -greedy policy π' with respect to Q_π is an improvement, then $V_{\pi'}(s) \geq V_\pi(s)$

$$\begin{aligned} Q_\pi(s, \pi'(s)) &= \sum_{a \in \mathcal{A}} \pi'(a|s) Q_\pi(s, a) \\ &= \frac{\epsilon}{m} \sum_{a \in \mathcal{A}} Q_\pi(s, a) + (1 - \epsilon) \max_{a \in \mathcal{A}} Q_\pi(s, a) \\ &\geq \frac{\epsilon}{m} \sum_{a \in \mathcal{A}} Q_\pi(s, a) + (1 - \epsilon) \sum_{a \in \mathcal{A}} \frac{\pi(a|s) - \frac{\epsilon}{m}}{1 - \epsilon} Q_\pi(s, a) \\ &= \sum_{a \in \mathcal{A}} \pi(a|s) Q_\pi(s, a) = V_\pi(s) \end{aligned}$$

Where $m = |\mathcal{A}|$.

Therefore from policy improvement theorem, $V_{\pi'}(s) \geq V_\pi(s)$

Monte-Carlo Policy Iteration



- Policy evaluation: Monte-Carlo policy evaluation. $Q_\pi(s, a)$
- Policy improvement: ϵ -greedy improvement

Is it necessary to wait until the end of execution of all trajectories ?

The mean μ_1, μ_2, \dots of a sequence x_1, x_2, \dots can be computed incrementally,

$$\begin{aligned}\mu_k &= \frac{1}{k} \sum_{j=1}^k x_j \\ &= \frac{1}{k} \left(x_k + \sum_{j=1}^{k-1} x_j \right) \\ &= \frac{1}{k} (x_k + (k-1)\mu_{k-1}) \\ &= \mu_{k-1} + \frac{1}{k}(x_k - \mu_{k-1})\end{aligned}$$

Incremental Monte-Carlo Updates

- Update $V(s)$ incrementally after episode $S_1, A_1, R_2, \dots, S_T$
- For each state S_t with return G_t :

$$N(S_t) \leftarrow N(S_t) + 1$$


$$V(S_t) \leftarrow V(S_t) + \frac{1}{N(S_t)}(G_t - V(S_t))$$

- In non-stationary problems, it can be useful to track a running mean, i.e. forget old episodes.

$$V(S_t) \leftarrow V(S_t) + \alpha(G_t - V(S_t))$$

How do we set α a.k.a **the learning rate** ?

Monte-Carlo Control



Every Iteration:

- Policy evaluation: Monte-Carlo policy evaluation. $Q \approx q_\pi$
- Policy improvement: ϵ -greedy improvement

On-policy first-visit MC control (for ε -soft policies), estimates $\pi \approx \pi_*$

Algorithm parameter: small $\varepsilon > 0$

Initialize:

$\pi \leftarrow$ an arbitrary ε -soft policy

$Q(s, a) \in \mathbb{R}$ (arbitrarily), for all $s \in \mathcal{S}$, $a \in \mathcal{A}(s)$

$Returns(s, a) \leftarrow$ empty list, for all $s \in \mathcal{S}$, $a \in \mathcal{A}(s)$

Repeat forever (for each episode):

Generate an episode following π : $S_0, A_0, R_1, \dots, S_{T-1}, A_{T-1}, R_T$

$G \leftarrow 0$

Loop for each step of episode, $t = T-1, T-2, \dots, 0$:

$G \leftarrow G + R_{t+1}$

Unless the pair S_t, A_t appears in $S_0, A_0, S_1, A_1, \dots, S_{t-1}, A_{t-1}$:

Append G to $Returns(S_t, A_t)$

$Q(S_t, A_t) \leftarrow \text{average}(Returns(S_t, A_t))$

$A^* \leftarrow \arg\max_a Q(S_t, a)$ (with ties broken arbitrarily)

For all $a \in \mathcal{A}(S_t)$:

$$\pi(a|S_t) \leftarrow \begin{cases} 1 - \varepsilon + \varepsilon/|\mathcal{A}(S_t)| & \text{if } a = A^* \\ \varepsilon/|\mathcal{A}(S_t)| & \text{if } a \neq A^* \end{cases}$$

Will this converge to an optimal policy?

How do we make sure exploration eventually stops?

Reinforcement
Learning

(SDMRL)

Sarah Keren

Recap

Model Free RL:
Monte Carlo

Temporal-
Difference
Learning

Comparing Monte
Carlo and TD
methods

Policy Search
Methods

How do we make sure exploration eventually stops?

- Decay ϵ Over Time: e.g., exponential decay $\epsilon_t = \gamma \cdot \epsilon_{t-1}$ with $\gamma \in (0, 1)$.
- After a sufficient number of iterations (when $Q(s, a)$ has stabilized), switch to a purely greedy policy.

Greedy in the Limit with Infinite Exploration (GLIE)

- All state-action pairs are explored infinitely many times,

$$\lim_{k \rightarrow \infty} N_k(s, a) = \infty$$

- The policy converges on a greedy policy,

$$\lim_{k \rightarrow \infty} \pi_k(a|s) = \mathbf{1}(a = \arg \max_{a' \in \mathcal{A}} Q_k(s, a'))$$

For example, ϵ -greedy is GLIE if α reduces to zero at $\epsilon_k = \frac{1}{k}$.

- Sample k th episode using $\pi : \{S_1, A_1, R_2, \dots, S_T\} \sim \pi$
- For each state S_t and action A_t in the episode,

$$N(S_t, A_t) \leftarrow N(S_t, A_t) + 1$$

$$Q(S_t, A_t) \leftarrow Q(S_t, A_t) + \frac{1}{N(S_t, A_t)}(G_t - Q(S_t, A_t))$$

- Improve policy based on new action-value function

$$\epsilon \leftarrow \frac{1}{k}$$

$$\pi \leftarrow \epsilon\text{-greedy}(Q)$$

Theorem

GLIE Monte-Carlo control converges to the optimal action-value function, $Q(s, a) \rightarrow Q^*(s, a)$

- Converge very slowly to correct utilities values (requires a lot of sequences)

$$V_{\pi}(s) = \mathbb{E}_{\pi} [G_t \mid S_t = s]$$

- Doesn't exploit Bellman constraints on policy values

$$V_{\pi}(s) = R(s) + \gamma \sum_{s'} \mathcal{P}(s'|s, \pi(s)) \cdot V_{\pi}(s')$$


- Can consider estimates that violate this property badly
- How can we incorporate such constraints ?

Temporal-Difference Learning

AB Example



Two states A, B; no discounting; 8 episodes of experience

- A, 0, B, 0
- B, 1
- B, 1
- B, 1
- B, 1
- B, 1
- B, 1
- B, 0



What is $V(A)$, $V(B)$?

From Monte Carlo to Temporal-Difference (TD) Learning



Reinforcement
Learning

(SDMRL)

Sarah Keren

Recap

Model Free RL:
Monte Carlo

Temporal-
Difference
Learning

Comparing Monte
Carlo and TD
methods

Policy Search
Methods

Reminder: Bellman Formulas

Bellman equation (for a given deterministic policy):

Multiple variations, e.g.,

$$V_{\pi}(s) = R(s) + \gamma \sum_{s'} \mathcal{P}(s'|s, \pi(s)) \cdot V_{\pi}(s')$$

Bellman equation (for a given stochastic policy):

$$V_{\pi}(s) = \sum_a \pi(a | s) \left[R(s, a) + \gamma \sum_{s'} \mathcal{P}(s' | s, a) V_{\pi}(s') \right]$$

Using Q

$$Q_{\pi}(s, a) = R(s, a) + \gamma \sum_{s'} \mathcal{P}(s' | s, a) \sum_{a'} \pi(a' | s') Q_{\pi}(s', a')$$

Can we use these in RL?

Bellman Optimality Equations

Reinforcement
Learning

(SDMRL)

Sarah Keren

Recap

Model Free RL:
Monte Carlo

Temporal-
Difference
Learning

Comparing Monte
Carlo and TD
methods

Policy Search
Methods

$$V^*(s) = \max_a \left[R(s, a) + \gamma \sum_{s'} \mathcal{P}(s' | s, a) V^*(s') \right]$$

$$Q^*(s, a) = R(s, a) + \gamma \sum_{s'} \mathcal{P}(s' | s, a) \max_{a'} Q^*(s', a')$$

$$V^*(s) = \max_a Q^*(s, a)$$

Can we use these in RL?

Reminder: during the learning process we see (possibly partial) trajectories of the form:

$$\tau = s_0, a_1, r_1, s_1, \dots, s_{T-1}, a_{T-1}, r_T, s_T$$

Temporal-Difference (TD) Learning: Evaluation

- For each transition from s to s' , we perform the following update

$$V_{\pi}(s) := V_{\pi}(s) + \alpha(R(s) + \gamma V_{\pi}(s') - V_{\pi}(s))$$

with α as the learning rate

How does this move us closer to satisfying the Bellman constraint ?

$$V_{\pi}(s) = R(s) + \gamma \sum_{s'} \mathcal{P}(s'|s, \pi(s)) \cdot V_{\pi}(s')$$

Reinforcement
Learning

(SDMRL)

Sarah Keren

Recap

Model Free RL:
Monte Carlo

Temporal-
Difference
Learning

Comparing Monte
Carlo and TD
methods

Policy Search
Methods

Temporal-Difference (TD) Learning: Evaluation

- For each transition from s to s' , we perform the following update

$$V_{\pi}(s) := V_{\pi}(s) + \alpha(R(s) + \gamma V_{\pi}(s') - V_{\pi}(s))$$

with α as the learning rate

How does this move us closer to satisfying the Bellman constraint ?

$$V_{\pi}(s) = R(s) + \gamma \sum_{s'} \mathcal{P}(s'|s, \pi(s)) \cdot V_{\pi}(s')$$

- $R(s) + \gamma V_{\pi}(s')$ is a (noisy) sample of utility based on the next state.
- The update maintains a “mean” of (noisy) utility samples
- How do we guarantee convergence to the true values ?

Temporal-Difference (TD) Learning: Evaluation

- For each transition from s to s' , we perform the following update

$$V_{\pi}(s) := V_{\pi}(s) + \alpha(R(s) + \gamma V_{\pi}(s') - V_{\pi}(s))$$

with α as the learning rate

How does this move us closer to satisfying the Bellman constraint ?

$$V_{\pi}(s) = R(s) + \gamma \sum_{s'} \mathcal{P}(s'|s, \pi(s)) \cdot V_{\pi}(s')$$

- $R(s) + \gamma V_{\pi}(s')$ is a (noisy) sample of utility based on the next state.
- The update maintains a “mean” of (noisy) utility samples
- How do we guarantee convergence to the true values ?
 - If the learning rate decreases appropriately with the number of samples (e.g. $\frac{1}{n}$), then the utility estimates will converge to true values (non-trivial).

Temporal-Difference (TD) Learning: Control

Reinforcement
Learning

(SDMRL)

Sarah Keren

Recap

Model Free RL:
Monte Carlo

Temporal-
Difference
Learning

Comparing Monte
Carlo and TD
methods

Policy Search
Methods

- Simplest temporal-difference learning algorithm: $TD(0)$
 - Update value $V(s_t)$ toward estimated return $R_{t+1} + \gamma V(S_{t+1})$

$$V(S_t) \leftarrow V(S_t) + \alpha(R_{t+1} + \gamma V(S_{t+1}) - V(S_t))$$

- If we are using Q

$$Q(S_t, A_t) \leftarrow Q(S_t, A_t) + \alpha(R_{t+1} + \gamma Q(S_{t+1}, A_{t+1}) - Q(S_t, A_t))$$

Which policy do we use here?

Temporal-Difference (TD) Learning

Do local updates of utility/value function on a per-action basis.

Reinforcement
Learning

(SDMRL)

Sarah Keren

Recap

Model Free RL:
Monte Carlo

Temporal-
Difference
Learning

Comparing Monte
Carlo and TD
methods

Policy Search
Methods

Temporal-Difference (TD) Learning

Do local updates of utility/value function on a per-action basis.

- TD methods learn directly from episodes of experience

Reinforcement
Learning

(SDMRL)

Sarah Keren

Recap

Model Free RL:
Monte Carlo

Temporal-
Difference
Learning

Comparing Monte
Carlo and TD
methods

Policy Search
Methods

Temporal-Difference (TD) Learning

Reinforcement
Learning

(SDMRL)

Sarah Keren

Recap

Model Free RL:
Monte Carlo

Temporal-
Difference
Learning

Comparing Monte
Carlo and TD
methods

Policy Search
Methods

Do local updates of utility/value function on a per-action basis.

- TD methods learn directly from episodes of experience
- TD is model-free: no learning of MDP transitions / rewards (doesn't try to estimate entire transition function).

Temporal-Difference (TD) Learning

Reinforcement
Learning

(SDMRL)

Sarah Keren

Recap

Model Free RL:
Monte Carlo

Temporal-
Difference
Learning

Comparing Monte
Carlo and TD
methods

Policy Search
Methods

Do local updates of utility/value function on a per-action basis.

- TD methods learn directly from episodes of experience
- TD is model-free: no learning of MDP transitions / rewards (doesn't try to estimate entire transition function).
- TD learns from incomplete episodes, by **bootstrapping** - updates a guess towards a guess.

- Goal: learn and optimize value function online from experience

- Incremental every-visit Monte-Carlo

- Update value $V(s_t)$ toward actual return G_t

$$V(S_t) \leftarrow V(S_t) + \alpha(G_t - V(S_t))$$


- Simplest temporal-difference learning algorithm: $TD(0)$

- Update value $V(s_t)$ toward estimated return $R_{t+1} + \gamma V(S_{t+1})$

$$V(S_t) \leftarrow V(S_t) + \alpha(R_{t+1} + \gamma V(S_{t+1}) - V(S_t))$$

- $R + \gamma V(S_{t+1})$ is called the **TD target**
 - $R_{t+1} + \gamma V(S_{t+1}) - V(S_t)$ is called the **TD error**

Going Home Example



Reinforcement
Learning

(SDMRL)

Sarah Keren

Recap

Model Free RL:
Monte Carlo

Temporal-
Difference
Learning

Comparing Monte
Carlo and TD
methods

Policy Search
Methods

Going Home Example

Reinforcement
Learning

(SDMRL)

Sarah Keren


Recap

Model Free RL:
Monte Carlo

Temporal-
Difference
Learning

Comparing Monte
Carlo and TD
methods

Policy Search
Methods



- τ_1 = office, 20, cable, 15, train, 30, train-park, 5, 2ndary, 15, home
- τ_2 = office, 20, cable, 15, train, 60, train-park, 5, 2ndary, 15, home
- τ_3 = office, 5, taub-park, 10, highway, 60, 2ndary, 15, home
- τ_4 = office, 5, taub-park, 10, highway, 60, 2ndary, 35, home
- τ_5 = office, 5, taub-park, 10, highway, 60


Going Home Example

$$V_{\pi}(S_t) \leftarrow V_{\pi}(S_t) + \alpha(G_t - V(S_t))$$

$$G_t = R_{t+1} + \gamma R_{t+2} + \dots + \gamma^{T-1} R_T$$

$$V_{\pi}(S_t) \leftarrow V_{\pi}(S_t) + \alpha(R(S_t) + \gamma V_{\pi}(S_{t+1}) - V_{\pi}(S_t))$$

$$\gamma = 1$$



- τ_1 = office, 20, cable, 15, train, 30, train-park, 5, 2ndary, 15, home
- τ_2 = office, 20, cable, 15, train, 60, train-park, 5, 2ndary, 15, home
- τ_3 = office, 5, taub-park, 10, highway, 60, 2ndary, 15, home
- τ_4 = office, 5, taub-park, 10, highway, 60, 2ndary, 35, home
- τ_5 = office, 5, taub-park, 10, highway, 60

Going Home Example

$$V_{\pi}(S_t) \leftarrow V_{\pi}(S_t) + \alpha(G_t - V_{\pi}(S_t))$$

$$G_t = R_{t+1} + \gamma R_{t+2} + \dots + \gamma^{T-1} R_T$$

$$V_{\pi}(S_t) \leftarrow V_{\pi}(S_t) + \alpha(R(S_t) + \gamma V_{\pi}(S_{t+1}) - V_{\pi}(S_t))$$

$$\gamma = 1$$

Exit office


Enter Cable

- τ_1 = office, 20, cable, 15, train, 30, train-park, 5, 2ndary, 15, home
- τ_2 = office, 20, cable, 15, train, 60, train-park, 5, 2ndary, 15, home
- τ_3 = office, 5, taub-park, 10, highway, 60, 2ndary, 15, home
- τ_4 = office, 5, taub-park, 10, highway, 60, 2ndary, 35, home
- τ_5 = office, 5, taub-park, 10, highway, 60

- Temporal-difference (TD) learning has several advantages over Monte-Carlo (MC)


- Temporal-difference (TD) learning has several advantages over Monte-Carlo (MC)
 - Lower variance
 - Online
 - Works with incomplete sequences
- Natural idea: use TD instead of MC in our control loop
 - Apply TD to $Q(S, A)$
 - Use ϵ -greedy policy improvement
 - Update every time-step.

Updating Action-Value Functions with



$$Q(S, A) \leftarrow Q(S, A) + \alpha(R + \gamma Q(S', A') - Q(S, A))$$

On-Policy Control With SARSA



Every time-step:

- Policy evaluation: Sarsa $Q \approx q_{\pi}$
- Policy improvement: ϵ -greedy improvement

Initialize $Q(s, a), \forall s \in \mathcal{S}, a \in \mathcal{A}(s)$, arbitrarily, and $Q(\text{terminal-state}, \cdot) = 0$

Repeat (for each episode):

 Initialize S

 Choose A from S using policy derived from Q (e.g., ϵ -greedy)

 Repeat (for each step of episode):

 Take action A , observe R, S'

 Choose A' from S' using policy derived from Q (e.g., ϵ -greedy)

$Q(S, A) \leftarrow Q(S, A) + \alpha[R + \gamma Q(S', A') - Q(S, A)]$


$S \leftarrow S'; A \leftarrow A';$

 until S is terminal

$$Q(S, A) \leftarrow Q(S, A) + \alpha(R + \gamma Q(S', A') - Q(S, A))$$

Additional topics:

- N-Step Sarsa
- Sarsa(λ)
- Forward and backward view Sarsa
- Convergence proof



- Evaluate target policy $\pi(a|s)$ to compute $\mathcal{V}_\pi(s)$ or $Q_\pi(s, a)$ while following behaviour policy $\mu(a|s)$


$$\{S_1, A_1, R_2, \dots, S_T\} \sim \mu$$

- Why is this important?
 - Learn from observing humans or other agents
 - Re-use experience generated from old policies $\pi_1, \pi_2, \dots, \pi_{t1}$
 - Learn about optimal policy while following exploratory policy
 - Learn about multiple policies while following one policy

Q-Learning

- We now consider off-policy learning of action-values $Q(s, a)$
- Next action is chosen using the behaviour policy
 $A_{t+1} \sim \mu(|S_t)$ but we update the target policy by considering the possible actions that could be applied to the next state and choose one with maximal value.

$$Q(S_t, A_t) \leftarrow Q(S_t, A_t) + \alpha(R_{t+1} + \gamma \max_{a'} Q(S', A') - Q(S_t, A_t))$$



Off-Policy Control with Q-Learning

- In Q-learning we allow both behaviour and target policies to improve
- The target policy π is greedy w.r.t. $Q(s, a)$

$$\pi(S_{t+1}) = \arg \max_{a'} Q(S_{t+1}, a')$$

- The behaviour policy μ is (e.g.) ϵ -greedy w.r.t. $Q(s, a)$
- The Q-learning target then simplifies:

$$\begin{aligned} & R_{t+1} + \gamma Q(S_{t+1}, A') \\ &= R_{t+1} + \gamma Q(S_{t+1}, \arg \max_{a'} Q(S_{t+1}, a')) \\ &= R_{t+1} + \max_{a'} \gamma Q(S_{t+1}, a') \end{aligned}$$

Reinforcement
Learning

(SDMRL)

Sarah Keren

Recap

Model Free RL:
Monte Carlo

Temporal-
Difference
Learning


Comparing Monte
Carlo and TD
methods

Policy Search
Methods

```
Initialize  $Q(s, a), \forall s \in \mathcal{S}, a \in \mathcal{A}(s)$ , arbitrarily, and  $Q(\text{terminal-state}, \cdot) = 0$ 
Repeat (for each episode):
  Initialize  $S$ 
  Repeat (for each step of episode):
    Choose  $A$  from  $S$  using policy derived from  $Q$  (e.g.,  $\epsilon$ -greedy)
    Take action  $A$ , observe  $R, S'$ 
     $Q(S, A) \leftarrow Q(S, A) + \alpha [R + \gamma \max_a Q(S', a) - Q(S, A)]$ 
     $S \leftarrow S'$ ;
  until  $S$  is terminal
```

$$Q(S_t, A_t) \leftarrow Q(S_t, A_t) + \alpha (R_{t+1} + \gamma \max_{a'} Q(S', A') - Q(S_t, A_t))$$

$$Q(S_t, A_t) \leftarrow Q(S_t, A_t) + \alpha(R_{t+1} + \gamma \max_{a'} Q(S', A') - Q(S_t, A_t))$$



Theorem

Q-learning control converges to the optimal action-value function, $Q(s, a) \rightarrow Q^*(s, a)$

SARSA vs. Q-Learning

Reinforcement
Learning

(SDMRL)

Sarah Keren

Recap

Model Free RL:
Monte Carlo

Temporal-
Difference
Learning

Comparing Monte
Carlo and TD
methods

Policy Search
Methods

SARSA:

$$Q(S_t, A_t) \leftarrow Q(S_t, A_t) + \alpha(R_{t+1} + \gamma Q(S_{t+1}, A') - Q(S_t, A_t))$$

Q-Learning:

$$Q(S_t, A_t) \leftarrow Q(S_t, A_t) + \alpha(R_{t+1} + \gamma \max_{a'} Q(S', A') - Q(S_t, A_t))$$

SARSA vs. Q-Learning

Initialize $Q(s, a), \forall s \in \mathcal{S}, a \in \mathcal{A}(s)$, arbitrarily, and $Q(\text{terminal-state}, \cdot) = 0$

Repeat (for each episode):

Initialize S

Choose A from S using policy derived from Q (e.g., ϵ -greedy)

Repeat (for each step of episode):

Take action A , observe R, S'

Choose A' from S' using policy derived from Q (e.g., ϵ -greedy)

$Q(S, A) \leftarrow Q(S, A) + \alpha [R + \gamma Q(S', A') - Q(S, A)]$

$S \leftarrow S'; A \leftarrow A';$

until S is terminal

Initialize $Q(s, a), \forall s \in \mathcal{S}, a \in \mathcal{A}(s)$, arbitrarily, and $Q(\text{terminal-state}, \cdot) = 0$

Repeat (for each episode):

Initialize S

Repeat (for each step of episode):


Choose A from S using policy derived from Q (e.g., ϵ -greedy)

Take action A , observe R, S'

$Q(S, A) \leftarrow Q(S, A) + \alpha [R + \gamma \max_a Q(S', a) - Q(S, A)]$

$S \leftarrow S';$

until S is terminal



Reinforcement
Learning

(SDMRL)

Sarah Keren

Recap

Model Free RL:
Monte Carlo

Temporal-
Difference
Learning

Comparing Monte
Carlo and TD
methods

Policy Search
Methods

Comparing Monte Carlo and TD methods

- TD can learn before knowing the final outcome
 - TD can learn online after every step
 - MC must wait until end of episode before return is known
- TD can learn without the final outcome
 - TD can learn from incomplete sequences
 - MC can only learn from complete sequences
 - TD works in continuing (non-terminating) environments
 - MC only works for episodic (terminating) environments

- Return $G_t = R_{t+1} + \gamma R_{t+2} + \dots + \gamma^{T-1} R_T$ is an **unbiased** estimate of $V_\pi(S_t)$
- True TD target $R_{t+1} + \gamma v_\pi(S_{t+1})$ is **unbiased estimate** of $\pi(S_t)$.
- TD target $R_{t+1} + \gamma V(S_{t+1})$ is a **biased** estimate $\pi(S_t)$. **why?**
- TD target has much lower variance than the return - **why?**

- Return $G_t = R_{t+1} + \gamma R_{t+2} + \dots + \gamma^{T-1} R_T$ is an **unbiased** estimate of $V_\pi(S_t)$
- True TD target $R_{t+1} + \gamma v_\pi(S_{t+1})$ is **unbiased estimate** of $\pi(S_t)$.
- TD target $R_{t+1} + \gamma V(S_{t+1})$ is a **biased** estimate $\pi(S_t)$. **why?**
- TD target has much lower variance than the return - **why?**
 - Return depends on many random actions, transitions, rewards
 - TD target depends on one random action, transition, reward

MC vs. TD (continued)

- MC has high variance, zero bias
 - Good convergence properties
 - (even with function approximation)
 - Not very sensitive to initial value
 - Very simple to understand and use
- TD has low variance, some bias
 - Usually more efficient than MC
 - TD(0) converges to $\pi(s)$
 - (but not always with function approximation)
 - More sensitive to initial value

Reinforcement
Learning

(SDMRL)

Sarah Keren

Recap



Model Free RL:
Monte Carlo

Temporal-
Difference
Learning

Comparing Monte
Carlo and TD
methods

Policy Search
Methods

Random Walk Example



Reinforcement
Learning

(SDMRL)

Sarah Keren

Recap

Model Free RL:
Monte Carlo

Temporal-
Difference
Learning

Comparing Monte
Carlo and TD
methods

Policy Search
Methods

- MC converges to solution with minimum mean-squared error
 - Best fit to the observed returns

$$\sum_{k=1}^K \sum_{t=1}^{T_k} (G_t^k - \mathcal{V}(s_t^k))^2$$

- $TD(0)$ converges to solution of max likelihood Markov model
 - Solution to the MDP $\langle \mathcal{S}, \mathcal{A}, \mathcal{P}, \mathcal{R}, \gamma \rangle$ that best fits the data

$$\mathcal{P}_{s,s'}^a = \frac{1}{N(s,a)} \sum_{k=1}^K \sum_{t=1}^{T_k} 1(s_t^k, a_t^k, s_{t+1}^k = s, a, s')$$


$$\mathcal{R}_s^a = \frac{1}{N(s,a)} \sum_{k=1}^K \sum_{t=1}^{T_k} 1(s_t^k, a_t^k = s, a) r_t^k$$

- TD exploits Markov property
 - Usually more efficient in Markov environments
- MC does not exploit Markov property
 - Usually more effective in non-Markov environments

Monte-Carlo vs. Temporal-Difference vs. Dynamic Programming Backup



Which is which ?

$$\mathcal{V}(s_t) \leftarrow \mathbb{E}_{\pi}[R_{t+1} + \gamma \mathcal{V}(S_{t+1})]$$



$$\mathcal{V}(S_t) \leftarrow \mathcal{V}(S_t) + \alpha(G_t \mathcal{V}(S_t))$$

$$\mathcal{V}(s_t) \leftarrow \mathcal{V}(s_t) + \alpha(R_{t+1} + \gamma \mathcal{V}(S_{t+1}) - \mathcal{V})]$$



(SDMRL)

Sarah Keren

Recap

Model Free RL:
Monte Carlo

Temporal-Difference
Learning

Comparing Monte
Carlo and TD
methods

Policy Search
Methods

Bootstrapping and Sampling

- **Bootstrapping:**

- MC
- DP
- TD

- **Sampling:**

- MC
- DP
- TD

Reinforcement
Learning

(SDMRL)

Sarah Keren

Recap

Model Free RL:
Monte Carlo

Temporal-
Difference
Learning

Comparing Monte
Carlo and TD
methods

Policy Search
Methods

Bootstrapping and Sampling

- **Bootstrapping:**
 - MC does not bootstrap
 - DP bootstraps
 - TD bootstraps
- **Sampling:**
 - MC samples
 - DP does not sample
 - TD samples

Reinforcement
Learning

(SDMRL)

Sarah Keren

Recap

Model Free RL:
Monte Carlo

Temporal-
Difference
Learning

Comparing Monte
Carlo and TD
methods

Policy Search
Methods

Unified View of Reinforcement Learning

Reinforcement Learning

(SDMRL)

Sarah Keren

Recap


Model Free RL:
Monte Carlo

Temporal-Difference
Learning

Comparing Monte
Carlo and TD
methods

Policy Search
Methods

Unified View



Extensions and Additional Topics (which we won't cover)

- n-Step Prediction
- Averaging n-Step Returns ($TD(\lambda)$)
- Eligibility Traces and credit assignment:
 - Frequency heuristic: assign credit to most frequent states
 - Recency heuristic: assign credit to most recent states
- Forward vs. backward view

Reinforcement
Learning

(SDMRL)

Sarah Keren

Recap

Model Free RL:
Monte Carlo

Temporal-
Difference
Learning

Comparing Monte
Carlo and TD
methods

Policy Search
Methods

Policy Search Methods

Policy Search

- **Key Idea:** Keep Twiddling the policy as long as its performance improves, then stop.
- In some ways, policy search is the simplest of all methods





Image from Guided Policy Search by Levine and Koltun, 2013


How does this fit within our general structure ?



Motivation for Policy-based Reinforcement Learning

Challenges:

- Dimensionality:
 - High-dimensional continuous state and action space
 - Huge variety of tasks
- Real world environments:
 - High-costs of generating data
 - Noisy measurements
- Exploration:
 - Do not damage the robot
 - Need to generate smooth trajectories



From

<https://icml.cc/2015/tutorials/PolicySearch.pdf>

Reinforcement
Learning

(SDMRL)

Sarah Keren

Recap

Model Free RL:
Monte Carlo

Temporal-
Difference
Learning

Comparing Monte
Carlo and TD
methods

Policy Search
Methods

Which policy search methods have we already seen ?

Reinforcement
Learning

(SDMRL)

Sarah Keren

Recap

Model Free RL:
Monte Carlo

Temporal-
Difference
Learning


Comparing Monte
Carlo and TD
methods

Policy Search
Methods

Which policy search methods have we already seen ?

- **Policy Iteration:** for known MDPs, we start with a fixed policy and iteratively improve until we reach convergence.
- **Hill Climbing:** maximize (or minimize) a target function $f(\mathbf{x})$. At each iteration, adjust a single element in \mathbf{x} and determine whether the change improves the value of $f(\mathbf{x})$.

Reminder: RL Approaches



Reinforcement
Learning

(SDMRL)

Sarah Keren

Recap

Model Free RL:
Monte Carlo

Temporal-
Difference
Learning

Comparing Monte
Carlo and TD
methods

Policy Search
Methods

Reminder: Value-based Reinforcement Learning:

Reinforcement
Learning

(SDMRL)

Sarah Keren

Recap

Model Free RL:
Monte Carlo

Temporal-
Difference
Learning

Comparing Monte
Carlo and TD
methods


Policy Search
Methods

- Estimate value function: e.g.

$$Q(s, a) \leftarrow Q(s, a) + \alpha(R + \gamma Q(s', a') - Q(s, a))$$

- Global estimate for all reachable states
 - Hard to scale to high-dimensional space
 - Approximations might compromise policy quality.
- Estimate global policy: e.g. $\pi'(s) = \arg \max_{a \in A} Q(s, a)$
 - Greedy policy update for all states
 - Policy update might get unstable
- Explore the whole state space: e.g. $\pi(a|s) = \frac{\exp(Q(s, a))}{\sum_a \exp(Q(s, a'))}$
 - Uncorrelated exploration in each step
 - (Might damage a robot)

Policy Search



Reinforcement
Learning

(SDMRL)

Sarah Keren


Recap

Model Free RL:
Monte Carlo

Temporal-
Difference
Learning

Comparing Monte
Carlo and TD
methods

Policy Search
Methods



- Policy search methods directly optimize policy parameters without learning a value function
- Key advantage: Can handle continuous action spaces naturally
- Different from value-based methods like Q-learning

- Policies can be represented as:
 - Neural networks: $\pi_{\theta}(a|s)$
 - Linear functions: $\pi_{\theta}(a|s) = \sigma(\theta^T \phi(s))$
 - Gaussian policies: $\pi_{\theta}(a|s) = \mathcal{N}(\mu_{\theta}(s), \Sigma_{\theta}(s))$

- A policy π is a function that maps states to actions (or state-action pairs to probabilities).
- We approximate the value or action-value function using parameters Θ ,

$$V_{\theta}(s) \approx V_{\pi}(s)$$

$$Q_{\theta}(s, a) \approx Q_{\pi}(s, a)$$

- A policy was generated directly from the value function e.g. using ϵ -greedy.
- **Policy search (gradient)** methods directly parametrise the policy $\pi_{\theta}(s, a) = \mathcal{P}[a|s, \theta]$
- Both model-free and model-based versions.

Policy Search Methods ¹

- Use parametrized policy $a \sim \pi(a|s; \theta)$, θ - parameter vector.
 - Compact parametrizations for high-dimensional spaces
 - Encode prior knowledge
- Locally optimal solutions e.g., $\theta_{new} = \theta_{old} + \alpha \frac{\partial J_{\theta}}{\partial \theta}$, where J is the loss function.
 - Safe policy updates
 - No global value function estimation
- Correlated local exploration
e.g.: $\theta \sim \mathcal{N}(\mu_{\theta}, \sigma_{\theta})$
 - Explore in parameter space
 - Generates smooth trajectories

<https://icml.cc/2015/tutorials/PolicySearch.pdf>

¹see Deisenroth, Neumann and Peters for A Survey of Policy Search for Robotics, FNT 2013

REINFORCE: Monte-Carlo Policy-Gradient Control (episodic) for π_*

Input: a differentiable policy parameterization $\pi(a|s, \theta)$

Algorithm parameter: step size $\alpha > 0$

Initialize policy parameter $\theta \in \mathbb{R}^{d'}$ (e.g., to $\mathbf{0}$)

Loop forever (for each episode):

 Generate an episode $S_0, A_0, R_1, \dots, S_{T-1}, A_{T-1}, R_T$, following $\pi(\cdot|\cdot, \theta)$

 Loop for each step of the episode $t = 0, 1, \dots, T - 1$:

$$G \leftarrow \sum_{k=t+1}^T \gamma^{k-t-1} R_k \quad (G_t)$$

$$\theta \leftarrow \theta + \alpha \gamma^t G \nabla \ln \pi(A_t | S_t, \theta)$$

Figure 2: Taken from Sutton and Barto

Recap and what next

- Spectrum of approaches to model-free RL
- Next: Dealing with large states space and model-based RL

